



SCHOOL OF ARCHITECTURE, COMPUTING AND ENGINEERING

Department of Engineering and Computing

Breaking the Tabular Ceiling: Heterogeneous Graph-Aware Multimodal Fusion for M&A Synergy Prediction

Hard Joshi

2512658

A report submitted in part fulfilment of the degree of
BSc (Hons) in *Data Science and Artificial Intelligence*

Supervisor: Arish Siddiqui
CN6000

08 May 2026

Abstract

Mergers and acquisitions destroy shareholder value in 70–90% of cases, yet existing predictive models consistently fail to identify value-destroying deals before announcement. This study argues that the root cause is architectural rather than computational: every prior generation of quantitative M&A model — from logistic regression to transformer-based NLP — treats each firm as an isolated data point, blind to the supply-chain topology, competitive network structure, and section-level semantic signals that collectively determine whether a proposed combination will create or destroy value. This dissertation addresses that gap by constructing, evaluating, and interpreting a heterogeneous multimodal architecture that encodes all three information dimensions simultaneously, and by formally testing whether each dimension carries predictive signal that is irreducible from the others.

A tri-modal late-fusion model (HeteroGraphSAGE) is built and evaluated on a universe of completed US domestic M&A transactions announced between 2000 and 2023, comprising approximately 3,000 deals with full multimodal coverage, sourced from Yahoo Finance, SEC EDGAR, and Bloomberg SPLC. Three feature blocks are fused via late concatenation: 56 financial ratio features (Block A), section-conditioned FinBERT embeddings of Management Discussion & Analysis and Risk Factor disclosures extracted separately from pre-announcement 10-K filings (Block B), and two-hop heterogeneous GraphSAGE embeddings derived from firm-level supply-chain networks (Block C). A dual-evaluation framework is applied: a classifier pipeline that predicts binary Cumulative Abnormal Return (CAR) direction and a regression pipeline that predicts CAR magnitude. Both pipelines are evaluated under five-fold stratified cross-validation with strict temporal splits and an 11-day event-window embargo to prevent forward-looking leakage. Three formal hypotheses are tested — H1 (Topological Alpha), H2 (Semantic Divergence), and H3 (Topological Arbitrage) — using paired t-tests, OLS regression, and Levene’s variance test respectively, with Bonferroni correction applied across all three tests.

All three hypotheses are supported. The full multimodal model achieves AUC-ROC = 0.5655, a statistically meaningful +0.0247 lift over the financial-only baseline (0.5408), confirming that supply-chain topology encodes directional signal that tabular models cannot recover (H1). OLS estimation on a semantic-divergence subsample of 1,140 deals recovers opposite-signed coefficients for MD&A similarity ($\beta = +0.0044$) and Risk Factor similarity ($\beta = -0.0080$), confirming that section-level semantic divergence is economically directional when the two sections are modelled separately rather than pooled (H2). Levene’s test across betweenness-centrality quantile groups yields $F = 7.07$ ($p = 0.0079$), confirming that structurally central acquirers experience statistically compressed announcement-return variance (H3). The study additionally establishes that adding undifferentiated FinBERT text to the financial baseline reduces AUC by -0.012 , demonstrating that naive NLP actively destroys predictive value when filing-section semantics are conflated — a methodological finding with direct implications for subsequent M&A NLP research. Continuous CAR magnitude remains intractable to regression across all configurations ($R^2 < 0$), precisely locating the boundary between tractable and structurally difficult M&A prediction sub-problems. All empirical results are rendered dynamically through the Deal Intelligence Terminal, an interactive research artefact ensuring full reproducibility.

Acknowledgments

This project would not have been possible without the support of several people who were generous with their time, guidance, and patience at various stages of a process that was significantly more ambitious than it looked on paper.

My deepest thanks go to my supervisor, Arish Siddiqui, whose guidance was direct, honest, and exactly what this project needed. He never let me settle for the convenient framing of a result or the easy conclusion, and the dissertation is considerably stronger for it. Having a supervisor who takes the work seriously is not something to take for granted.

I am grateful to the University of East London for providing access to institutional data infrastructure, Bloomberg Terminal access, and the computational environment that made the empirical programme feasible. Some of the datasets used in this project — supply-chain topology, institutional equity returns, and SEC filing archives — each individually exceed what I could have accessed independently. The fact that this research was conducted at Bachelors level is, in large part, a function of that access.

This project was built in the company of people who kept things in perspective when the results were stubborn and the deadline was not. Thank you to my friends and family — particularly those who sat through more than one explanation of why a negative R^2 is not a failure — for their interest, their humour, and their patience.

Contents

Abstract	i
Acknowledgments	ii
1 Introduction	1
1.1 The Problem: A Century of Predictable Failure	1
1.2 Aims and Objectives	1
1.3 The Proposed Solution	2
1.4 Research Hypotheses	2
1.5 What This Study Contributes	3
1.6 Structure of the Dissertation	3
2 Literature Review	4
2.1 The Central Argument: A Gap Three Decades Wide	4
2.2 Stream I: The M&A Paradox and the Failure of Valuation Theory	5
2.2.1 The Empirical Record: Why Deals Keep Failing	5
2.2.2 Behavioural Distortions: When Executive Confidence Contaminates the Data	5
2.2.3 Market Timing: When Deal Waves Are About Valuation, Not Strategy	6
2.2.4 The Epistemological Barrier: What Due Diligence Cannot See	6
2.3 Stream II: The Tabular Paradigm — An Asymptotic Ceiling	7
2.3.1 The Econometric Era: When Linearity Assumed Too Much	7
2.3.2 Why Machine Learning Could Not Fully Escape the Same Trap	7
2.3.3 Why Deep Neural Networks Did Not Break Through	8
2.4 Stream III: The Semantic Turn — Powerful Tools, Wrong Targets	9
2.4.1 The Bag-of-Words Era: A Useful Foundation with Hard Limits	9
2.4.2 The Transformer Revolution: A Better Tool, But Often Misdirected	10
2.4.3 Why Standard FinBERT Cannot Fully Recover the Textual Signal	10
2.5 Stream IV: The Topological Turn — Network Structure as Irreducible Signal	11
2.5.1 The Economic Foundation: Supply Chain Momentum	11
2.5.2 Corporate Network Centrality and Value	12
2.5.3 Why Adding Centrality to XGBoost Is Insufficient	13
2.5.4 Why Edge-Type Semantics Matter: The Case for Heterogeneous Graphs ..	13
2.6 The Measurement Foundation: Event Study Methodology	14
2.6.1 How Announcement Returns Are Measured	14
2.6.2 Event Window Choice	15
2.6.3 Why Binary Classification Rather Than Regression	15
2.7 The Multimodal Imperative: Why Fusion Is the Only Exit	16
2.7.1 Complementary Variance: Three Lenses on the Same Deal	16
2.7.2 Why HeteroGraphSAGE Fusion Is the Logical Response	16
2.7.3 Architectural Choice: Late Fusion and Its Rationale	18
2.8 Synthesis: Gap Table and Research Hypotheses	19
2.8.1 Formal Research Hypotheses	19
3 Methodology	20
3.1 Introduction	20
3.2 Research Philosophy and Design	21

3.3	Research Hypotheses	22
3.4	Data Sources and Collection	23
3.4.1	M&A Deal Universe	23
3.4.2	Equity Return Data	24
3.4.3	Textual Data (SEC EDGAR)	24
3.4.4	Supply-Chain Network Data	24
3.5	Data Preprocessing	24
3.5.1	Cleaning and Quality Control	24
3.5.2	Feature Engineering and Normalisation	24
3.5.3	Temporal Splitting and Event-Window Embargo	24
3.5.4	Missing Data Strategy	25
3.6	Feature Extraction	25
3.6.1	Block A — Financial Features	26
3.6.2	Block B — Textual Features (FinBERT)	27
3.6.3	Block C — Graph Features (HeteroGraphSAGE)	29
3.7	Model Architecture	32
3.7.1	Baseline Models	32
3.7.2	Fusion Model Architecture	33
3.7.3	Training Configuration	34
3.8	Target Variable: Cumulative Abnormal Return	35
3.8.1	Stage 1: OLS Market Model	35
3.8.2	Abnormal Returns and CAR	36
3.8.3	Stage 2: Model Prediction vs. Actual CAR	37
3.9	Hypothesis Testing	39
3.10	Evaluation Metrics	39
3.11	Ethical, Legal, and Social Considerations	39
4	Findings and Discussion	41
4.1	Introduction	41
4.2	Classification Results	41
4.2.1	The Classification Ablation Ladder	41
4.2.2	Hyperparameter Tuning: An Honest Negative Result	43
4.3	Regression Results	43
4.3.1	Continuous CAR Remains Structurally Difficult	43
4.4	The M2 Reversal	44
4.4.1	Why Naive NLP Makes the Model Worse	44
4.5	Hypothesis Tests	45
4.5.1	H1 — Topological Alpha (Supported)	45
4.5.2	H2 — Semantic Divergence (Supported)	46
4.5.3	H3 — Topological Arbitrage: Information Transparency Dampening (Supported)	47
4.6	Interpretability	48
4.6.1	SHAP and Economic Credibility	48
4.7	Practical Meaning	49
4.7.1	Why a +0.025 AUC Gain Matters	49

4.8	Limitations	50
4.8.1	Boundaries of the Findings	50
4.9	Synthesis	50
5	Evaluation	51
5.1	Introduction	51
5.2	Product Evaluation	51
5.2.1	Objectives and Achievement	51
5.2.2	Performance Against the Literature	52
5.2.3	Product Limitations	52
5.2.4	The Deal Intelligence Terminal as Product	53
5.3	Process Evaluation	54
5.3.1	Research Design Decisions That Worked	54
5.3.2	Research Design Decisions That Underperformed	54
5.3.3	Scope and Time Management	55
5.3.4	Reflections on the Research Question	56
5.4	Personal Reflection and Mental Wealth	56
6	Conclusion	58
6.1	The Core Argument Resolved	58
6.2	Synthesis of Empirical Findings	58
6.3	Research Contributions	59
6.4	Future Opportunities	59
6.5	Reflection on Aims, Objectives, and Personal Development	60
6.5.1	Achievement of Objectives and Project Trajectory	60
6.5.2	Difficulties and Learning Curve	60
6.5.3	Real-World Use Case and Transferable Skills	60
6.6	Final Remark	61
	Bibliography	62
A	Initial and Revised Project Aims and Objectives	67
A.1	Initial Project Aims and Objectives (As Proposed)	67
A.2	Evolved Aims and Objectives (Final Implementation)	68
A.3	Summary of Key Divergences	69
B	Final Project Proposal	70
B.1	Project Metadata	70
B.2	Project Title	70
B.3	Revised Aim	70
B.4	Research Hypotheses	70
B.5	Scope and Dataset	71
B.5.1	Deal Universe	71
B.5.2	Feature Blocks	71
B.5.3	Target Variable	72
B.5.4	Leakage Prevention	72
B.6	Architecture	72
B.6.1	Feature Extraction	72
B.6.2	Fusion Strategy	72

B.6.3	Evaluation Framework	72
B.6.4	Hypothesis Testing Protocol	73
B.7	Ablation Ladder	73
B.8	Research Output Artefact	73
B.9	Key Divergences from Initial Proposal	74
B.10	Ethical, Legal, and Social Considerations	74
C	Application for Approval of Research Activities	75
D	Deal Intelligence Terminal	81
D.1	The Model Evidence Wall	81
D.2	Deal-Level Diagnostics	83
D.2.1	The Synergy Forecast	83
D.2.2	Feature Interaction Zones	84
D.2.3	The Glass Box (SHAP Attribution)	85
D.3	Supplementary Research Modules	86
D.3.1	The Methodology Engine	86
D.3.2	The Evaluation Lab	87
D.3.3	The Hypothesis Lab	88
D.3.4	Pipeline Architecture	89
E	Selected Implementation Code Snippets	90
E.1	Selection Criteria	90
E.2	Two-Stage CAR Pipeline	91
E.3	Temporal Splitting and Embargo Control	92
E.4	Section-Aware FinBERT Extraction	93
E.5	Heterogeneous Graph Construction and Embedding	95
E.6	Late-Fusion Prediction Head	96
E.7	SHAP-Based Interpretation	97
E.8	Repository Availability	97

List of Figures

Figure 1	System architecture pipeline: Multimodal feature extraction, fusion, and evaluation.	21
Figure 2	Section-aware FinBERT NLP pipeline extracting MD&A and Risk Factors semantic vectors.	27
Figure 3	HeteroGraphSAGE embedding process for supply-chain topological representation.	29
Figure 4	Supply-chain graph construction showing inter-firm dependencies and topological topology.	30
Figure 5	Multimodal late-fusion architecture combining financial, textual, and graph embeddings.	33
Figure 6	ROC curves comparing financial-only baseline (M1) against multimodal fusion (M3). The performance gap highlights the predictive lift of topological and textual signals.	42
Figure 7	Supply-chain ego-network of a highly central acquirer firm. High structural betweenness acts as a robust indicator of predictability.	45
Figure 8	AUC performance by model variant across different economic sectors. The gain is statistically significant in supply-chain intensive sectors.	46
Figure 9	H2 coefficient direction showing opposing market reactions to MD&A versus Risk Factor textual similarity.	47
Figure 10	Volatility Funnel demonstrating variance compression in CAR as graph centrality increases (H3).	48
Figure 11	Global SHAP manifold for the M3 fusion model, highlighting the relative feature importance across tabular, text, and graph modalities. Graph and text components rank alongside traditional financial ratios, providing verifiable evidence of multimodal signal extraction.	49
Figure 12	Ethical approval form — page 1 of 5.	76
Figure 13	Ethical approval form — page 2 of 5.	77
Figure 14	Ethical approval form — page 3 of 5.	78
Figure 15	Ethical approval form — page 4 of 5.	79
Figure 16	Ethical approval form — page 5 of 5.	80
Figure 17	The Ablation Wall interface comparing the M1 (financial-only) baseline against the full M3 (financial + text + graph) architecture across AUC-ROC, F1, precision, and recall.	82
Figure 18	The diagnostic verdict showing a 91% synergy probability calibrated via the M3 ensemble, alongside betweenness centrality rank (30th percentile) and MD&A semantic match.	83
Figure 19	Zone 1 (Topological Embeddings) rendering the acquirer–target supply-chain subgraph, and Zone 2 (Semantic Radar) plotting the pentagonal linguistic profile comparison.	84
Figure 20	Zone 3: SHAP waterfall decomposition of the synergy probability, with the Algorithmic Translation summary below.	85
Figure 21	The Methodology Engine interface, visualising the end-to-end M3 pipeline as an interactive Directed Acyclic Graph (DAG) with stage-specific output statistics. .	86

Figure 22	The Evaluation Lab, providing a structured defence of AUC-ROC as the primary performance metric with an interactive ROC curve visualisation.	87
Figure 23	The Hypothesis Lab, showcasing the sector-wise AUC lift used to validate Hypothesis 1 (Topological Alpha) with $\pm 1\sigma$ error bars.	88
Figure 24	The Pipeline Architecture interface, documenting the internal tensor flow and the 249-dimensional multimodal feature concatenation logic.	89

List of Tables

Table 1	Methodological Evolution Matrix (Part I): Evolution of M&A prediction research through three paradigms.	17
Table 2	Methodological Evolution Matrix (Part II): The transition to topological and multimodal fusion.	18
Table 3	Research Hypotheses	23
Table 4	Temporal splits and embargo design	25
Table 5	Feature block definitions and fusion-pipeline construction steps.	26
Table 6	HeteroGraphSAGE architecture and training configuration.	31
Table 7	Model variants for ablation experiments.	32
Table 8	Event study timeline. The gap prevents estimation-period contamination of the CAR window.	36
Table 9	Evaluation metrics. Regression metrics (MAE, RMSE, R^2 , Huber) apply to the Regressor Pipeline; classification metrics (Dir. Accuracy, AUC-ROC, F1) apply to the Classifier Pipeline.	38
Table 10	Classification ablation ladder — best model result per feature configuration under five-fold stratified cross-validation. Bold indicates headline AUC result.	42
Table 11	Regression pipeline summary — representative continuous CAR results. Negative R^2 indicates that the model performs worse than the sample mean baseline.	43
Table 12	Contextualised product performance. Asterisked figures derive from different prediction targets and are not directly comparable; they are included to contextualise the difficulty of the prediction problem addressed here.	52
Table 13	Summary of methodological divergences between the initial proposal and final implementation.	69
Table 14	Project metadata as submitted for CN6000 approval.	70
Table 15	Summary of multimodal feature blocks.	71
Table 16	Formal hypothesis testing protocol.	73
Table 17	Ablation ladder configuration.	73
Table 18	Summary of project trajectory and methodological divergences.	74

Chapter 1: Introduction

1.1 The Problem: A Century of Predictable Failure

Every year, hundreds of companies announce mergers and acquisitions worth trillions of dollars. Each deal rests on the same fundamental bet: that the combined entity will generate more economic value than the two firms could produce independently. The gap between that bet and its typical outcome is one of the most durable findings in modern corporate finance. Between 70% and 90% of acquisitions fail to generate value for acquirer shareholders, a failure rate that has remained statistically stable across a century of takeover activity and through every wave of increasingly sophisticated due diligence technology (Martynova and Renneboog, 2008). The question this dissertation asks is not **why** individual deals fail — the behavioural and market-timing explanations for that are well-established. The question is structural: what information do all existing predictive frameworks systematically miss, and is that information recoverable?

The answer proposed here is that existing models fail not because of insufficient computational power, but because of a flawed representational assumption. Every major class of quantitative M&A predictor — from logistic regression on financial ratios to gradient-boosted trees and transformer-based NLP pipelines — treats each firm as an isolated data point. This means that no matter how sophisticated the algorithm, it operates in a feature space that is topologically flat: it sees the numbers on the balance sheet but cannot see the economic ecosystem the firm inhabits. A firm's largest customer may be on the verge of bankruptcy. Its target may share fragile single-source suppliers. The combined entity may create a bottleneck in a critical industrial network. None of these facts appear in a financial ratio vector, and none can be recovered by adding more layers to the same architecture.

1.2 Aims and Objectives

The central aim of this study is to construct, evaluate, and interpret a heterogeneous multimodal machine learning framework for M&A synergy prediction that transcends traditional tabular limitations by integrating financial, textual, and topological data.

The project was guided by the following core objectives:

1. **Literature Synthesis:** To conduct a comprehensive review of existing econometric, NLP, and network-based M&A prediction models, identifying the structural limitations of isolated analytical streams.
2. **Multimodal Dataset Construction:** To engineer a unified dataset combining historical M&A financial fundamentals, section-specific 10-K regulatory filings, and inter-firm supply-chain relationship graphs.
3. **Architectural Development:** To design and train a late-fusion neural architecture capable of processing this heterogeneous feature space.

4. **Benchmark Evaluation:** To rigorously evaluate the multimodal model’s predictive performance (AUC-ROC) against a standard financial-only baseline to quantify the incremental value of non-financial signals.
5. **Interpretability and Feature Attribution:** To utilise SHAP (SHapley Additive exPlanations) to deconstruct model predictions, identifying the precise economic drivers of synergy and interpreting the interaction between financial fundamentals and ecosystem topology.

1.3 The Proposed Solution

This dissertation proposes and empirically tests a multimodal late-fusion architecture designed to break through the tabular ceiling by encoding three complementary information modalities simultaneously:

- **Block A — Financial Fundamentals:** 56 ratio-level features drawn from Yahoo Finance, covering acquirer and target leverage, liquidity, profitability, and deal structure. This is the baseline information available to every prior quantitative model.
- **Block B — Section-Aware Textual Semantics:** FinBERT-encoded embeddings from the Management Discussion & Analysis (MD&A) and Risk Factors sections of pre-announcement 10-K filings, processed separately so that their opposing economic signals are not cancelled by aggregation.
- **Block C — Heterogeneous Graph Topology:** Node embeddings derived from a two-hop HeteroGraphSAGE model trained on Bloomberg SPLC supply-chain relationship data, encoding each firm’s structural position within the industrial network it actually inhabits.

The three blocks are fused via late concatenation and evaluated across a dual-pipeline framework: a classifier pipeline that predicts the **direction** of acquirer Cumulative Abnormal Return (CAR) at announcement, and a regression pipeline that attempts to predict CAR **magnitude**. As Chapter 4 demonstrates, these two pipelines resolve differently — and that difference is itself a substantive finding. To ensure absolute methodological rigour, the pipeline enforces strict temporal splitting and an 11-day event-window embargo to explicitly prevent target leakage and market microstructure contamination.

1.4 Research Hypotheses

The study tests three formal hypotheses derived from the four knowledge streams surveyed in Chapter 2. Each hypothesis corresponds to one of the three modality blocks and is designed to be falsifiable within the available dataset.

H1 — The Topological Alpha Hypothesis: The inclusion of heterogeneous supply-chain graph topology will yield a statistically significant improvement in directional deal discrimination (AUC-ROC) over a financial-only baseline, proving that topological data yields orthogonal predictive alpha.

H2 — The Semantic Divergence Hypothesis: The relationship between textual similarity and announcement returns is conditional on document section. MD&A similarity (strategic fit)

correlates positively with CAR; Risk Factor similarity (shared liability exposure) correlates negatively. These opposing effects are mathematically suppressed when the two sections are pooled.

H3 — The Topological Arbitrage Hypothesis: Acquirers with high betweenness centrality in the supply-chain graph exhibit statistically compressed variance in announcement return outcomes relative to peripheral acquirers, as measured by Levene’s test across centrality quantile groups.

1.5 What This Study Contributes

The academic contribution of this work lies at the intersection of three fields that have not previously been unified in the context of post-merger synergy prediction. Supply chain finance, financial NLP, and heterogeneous graph neural networks have each advanced substantially as independent research programmes. This dissertation is, to the best of current knowledge, the first study to direct their combination at binary CAR direction classification as a proxy for M&A synergy.

The contribution is not only architectural. By reporting the M2 reversal — the empirical finding that adding aggregated FinBERT text **degrades** prediction quality — the study provides a methodological warning to future researchers that NLP is not automatically useful in M&A prediction unless section semantics are preserved. By reporting negative R^2 values on the regression pipeline, the study defines the boundary between what multimodal architecture can and cannot achieve in this domain. And by building the Deal Intelligence Terminal, the study operationalises the theoretical framework into a live, reproducible research artefact that renders all empirical claims interactively.

1.6 Structure of the Dissertation

The dissertation proceeds as follows. Chapter 2 builds a sustained critical argument across four interconnected knowledge streams, demonstrating that each wave of prior scholarship produced an asymptotic ceiling and identifying precisely the structural information each paradigm discarded. Chapter 3 details the full research design, data pipeline, feature construction, and hypothesis-testing protocol. Chapter 4 reports the empirical findings across both pipelines, including the negative results that establish the boundary conditions of the architecture. Chapter 5 evaluates both the research product and the research process honestly, contextualising performance against the literature and identifying what should be done differently in future work. Chapter 6 synthesises the findings and places the contribution in the context of the broader financial machine learning research agenda.

Chapter 2: Literature Review

2.1 The Central Argument: A Gap Three Decades Wide

Every year, thousands of companies announce mergers and acquisitions — deals in which one firm pays a premium to absorb another, betting that the combined entity will be worth more than the sum of its parts. Yet decades of evidence suggest that this bet fails more often than it succeeds. The central question motivating this study is simple: if this failure rate has persisted across a century of increasingly sophisticated corporate due diligence, what information are existing analytical frameworks systematically missing?

This review builds a sustained argument across four interconnected knowledge streams, each representing a distinct wave of scholarship that has tried — and partially failed — to answer that question. Rather than a neutral survey, each stream is examined for what it structurally could not achieve, with each identified gap mapped directly to one of the three hypotheses tested in this dissertation.

1. **Stream I — The M&A Paradox and the Failure of Valuation Theory:** diagnoses why synergy prediction is genuinely difficult and establishes that the root cause is not a lack of computing power, but a structural information blindspot.
2. **Stream II — The Tabular Paradigm:** traces the full history of quantitative M&A models — from logistic regression through gradient-boosted trees and early deep neural networks — and demonstrates that every generation inherited the same core architectural limitation, producing an asymptotic accuracy ceiling no amount of algorithmic sophistication could break through.
3. **Stream III — The Semantic Turn:** examines how natural language processing entered the M&A prediction landscape and why, despite genuinely powerful tools, the field systematically pointed them at the wrong question — and even when pointed correctly, fell into a structural trap that cancelled the very richness it sought to exploit.
4. **Stream IV — The Topological Turn:** engages with supply chain finance, corporate network theory, and graph neural network research to establish that the structure of a firm’s external relationships — its suppliers, customers, and competitors — carries economic information about deal value that is, by mathematical necessity, invisible to any model that treats each firm as an isolated data point.

The review concludes that these four streams converge on a single architectural response: a heterogeneous graph model that simultaneously fuses financial fundamentals (Block A), section-conditioned textual embeddings (Block B), and supply-chain and competition graph topology (Block C). As of 2025, no published study has directed this combined architecture at post-merger synergy outcome prediction.

2.2 Stream I: The M&A Paradox and the Failure of Valuation Theory

2.2.1 The Empirical Record: Why Deals Keep Failing

Mergers and acquisitions are the primary mechanism through which capital is reallocated across the global economy. When a company acquires another, the underlying premise is that the combined entity will generate more value than the two firms could independently — through cost savings, revenue growth, market power, or access to new capabilities. This theoretical surplus is called “synergy.”

The empirical record, however, tells a more troubling story. (Martynova and Renneboog, 2008) documented that between 70% and 90% of acquisitions fail to generate value for the acquirer’s shareholders — a finding replicated across multiple independent research programmes (Christensen *et al.*, 2011). More striking than the magnitude of this failure rate is its persistence: it has remained statistically stable across entirely different economic eras, from the conglomerate wave of the 1960s through the technology boom of the late 1990s and the post-financial-crisis consolidation wave of the 2010s (Martynova and Renneboog, 2008). A failure rate that survives such varied conditions is unlikely to be cyclical; it points toward something structurally wrong with how deals are evaluated.

(Bradley, Desai and Kim, 1988) established a revealing asymmetry in the data: while shareholders of acquired firms (the targets) captured average announcement-day abnormal returns of approximately 30%, the shareholders of acquiring firms systematically lost value, with combined deal synergy gains averaging only 7.4% of combined pre-deal firm value. The gap between what acquiring firms paid and what the market believed the combination was worth is precisely what this study attempts to predict.

2.2.2 Behavioural Distortions: When Executive Confidence Contaminates the Data

The first layer of explanation for persistent M&A failure is behavioural. (Roll, 1986) proposed the Hubris Hypothesis: acquiring managers systematically overestimate their own ability to extract value from a target, treating the market’s price for that target as an underestimate rather than as a rational, information-weighted signal. This overconfidence drives them to pay premiums that the actual post-merger performance cannot justify — what (Sirower, 1997) termed the “Synergy Trap.”

From a machine learning perspective, this creates a specific and irreducible data quality problem. If the acquisition premiums recorded in deal databases partly reflect executive overconfidence rather than genuine synergy potential, then any model trained to predict outcomes from those premiums is fitting a signal that has been systematically contaminated with human behavioural error. No amount of regularisation or hyperparameter tuning can remove bias that is baked into the target variable itself. This is one of the reasons this study uses Cumulative Abnormal Return (CAR) — the stock market’s immediate revision of its estimate of firm value upon hearing the deal announcement — as the prediction target, rather than deal premia or accounting-based synergy measures. CAR represents the market’s collective, forward-looking

assessment of deal value and is considerably less susceptible to the narrative distortions of individual executives (MacKinlay, 1997).

2.2.3 Market Timing: When Deal Waves Are About Valuation, Not Strategy

A complementary source of noise in M&A data was identified by (Shleifer and Vishny, 2003). Their stock-market-driven acquisitions model demonstrates that companies with temporarily overvalued equity have a rational incentive to use their inflated shares as acquisition currency — essentially exchanging overpriced paper for real assets. Under this model, a substantial fraction of any M&A wave is driven by relative misvaluation rather than genuine operational complementarity between the two firms. The observed announcement-period returns, in these cases, reflect markets correcting for overvaluation rather than pricing synergy realisation. (Grossman and Stiglitz, 1980) established the theoretical boundary here: in a world where gathering and processing information is costly, markets cannot be perfectly efficient, and persistent information asymmetries around deal announcements are a natural and stable equilibrium. This bounded-efficiency framing is what makes the prediction problem tractable in the first place — if markets perfectly and instantaneously priced all available information, there would be no learnable signal left.

2.2.4 The Epistemological Barrier: What Due Diligence Cannot See

Even setting aside behavioural and timing distortions, the fundamental problem is that conventional due diligence is architecturally blind to certain categories of risk. (Angwin, 2001) documented that standard due diligence operates in organisational silos: financial, legal, and operational teams independently assess their respective domains without modelling how those domains interact. This fragmentation means that multiplicative risks are systematically missed. Consider a firm with strong cash flows that happens to source 80% of its critical components from a single supplier. Each dimension passes a standard due diligence check in isolation; the conjunction — that the entire cash-flow advantage is conditional on a single supply relationship — is invisible to teams that never share information.

(Akerlof, 1970) provided the theoretical foundation for why mandated disclosures cannot fully resolve this problem. His “Market for Lemons” framework shows that when buyers and sellers have asymmetric information, sellers have an incentive to obscure their vulnerabilities within the noise of standard reporting formats. A 10-K filing, however detailed, is ultimately a document a company prepares about itself. What it cannot obscure — at least not easily — is the verifiable structure of its external relationships: which firms it buys from, sells to, and competes with. This network of real economic ties is precisely what Block C in this study’s architecture encodes, and it is the one source of synergy-relevant information that is most resistant to strategic misrepresentation.

Stream I conclusion: The M&A failure rate is partly predictable from structural information that existing valuation frameworks systematically exclude. The following three streams each diagnose a specific category of excluded information and the architectural limitations that kept it out.

2.3 Stream II: The Tabular Paradigm — An Asymptotic Ceiling

2.3.1 The Econometric Era: When Linearity Assumed Too Much

Before examining the historical progression of quantitative M&A models, the structural limitation of the “tabular” paradigm must be defined. A tabular model takes a spreadsheet-style matrix of numbers — one row per deal, one column per financial ratio or deal characteristic — and learns statistical patterns that separate value-creating deals from value-destroying ones. Each deal is treated as an independent, self-contained data point. The model has no mechanism for knowing that two deals involve acquirers who share a supplier, or that the target in one deal is a customer of the acquirer in another. It sees the numbers; it does not see the relationships.

The first generation of quantitative M&A models — most influentially (Palepu, 1986) and (Barnes, 1990) — used logistic regression and multiple discriminant analysis on financial ratios. These approaches assume that each financial variable has a monotonic, linear relationship with the outcome: more of a given ratio always either increases or decreases the probability of a successful deal. (Barnes, 1990) demonstrated empirically that this assumption breaks down for leverage. Moderate debt signals fiscal discipline and managerial confidence; very high debt signals financial fragility. The relationship is non-monotonic — it changes direction at some threshold — and linear models fundamentally cannot represent direction changes.

More revealing still, (Palepu, 1986) own out-of-sample validation found that acquisition targets could not be predicted reliably beyond chance levels. This was not a data collection problem or an implementation flaw; it was evidence that the information contained in financial ratios was simply insufficient to reconstruct synergy outcomes. The consistently low pseudo- R^2 values (below 0.10) documented across econometric studies (Betton, Eckbo and Thorburn, 2008) are not calibration failures — they are a measurement of the information gap that financial ratios leave behind. This ceiling is the first empirical benchmark any successor model must demonstrably improve upon.

2.3.2 Why Machine Learning Could Not Fully Escape the Same Trap

The arrival of ensemble machine learning methods — random forests, gradient-boosted decision trees, and ultimately XGBoost — promised to resolve the non-linearity problem. (Zhang *et al.*, 2024) and similar studies deployed these techniques on financial ratio vectors and did report accuracy improvements over logistic baselines. Those improvements are real, but they are bounded by four structural problems that persist regardless of how sophisticated the algorithm becomes.

The Independence Ceiling: Both random forests and XGBoost operate on rows of a feature table where each row represents one deal in complete isolation. The algorithm learns patterns within individual deals but has no mechanism to propagate information between related deals. If two acquiring firms share a common supplier, and that supplier is experiencing financial distress, both deals carry a latent risk that is invisible in their individual financial ratios. This is not a data-engineering problem that can be solved by adding more features — it is an architectural

impossibility for any model that processes deals independently (Hamilton, Ying and Leskovec, 2017).

Survivorship Bias: M&A datasets are almost exclusively drawn from completed transactions. Deals that were considered but abandoned — often precisely because due diligence revealed serious problems — never appear in the training data. This systematically biases estimated model coefficients toward the characteristics of deals that proceeded, rather than deals that actually created value (Betton, Eckbo and Thorburn, 2008). No machine learning algorithm can correct for data that was never collected.

Look-Ahead Contamination: Several prominent ML studies, including (Zhang *et al.*, 2024), construct financial features from annual filings without explicitly verifying that the fiscal year end-date precedes the deal announcement date. When a model trains on financial data that postdates the deal it is predicting, its apparent accuracy reflects information that would not have been available at the time of the decision — inflating reported performance beyond any economically realisable level.

Crystallised Analyst Bias: A common preprocessing step in these pipelines is discretising continuous financial variables into categorical bins before feeding them into tree models. This encodes whatever assumptions the analyst made when deciding where to draw the bin boundaries into the feature representation in a way that cannot be undone. The model then learns from a filtered, pre-interpreted version of the data rather than the raw signal — an automated version of the same heuristic reasoning employed by human analysts whose systematic failures are the very motivation for building quantitative models in the first place.

Taken together, these four problems create what can be described as an asymptotic accuracy ceiling: once a tabular model has captured all the linear, rank-order, and non-parametric signal available in deal-level financial features, no further improvement is achievable by making the model more sophisticated. The ceiling is a function of the information content of the feature space, not of computational power.

2.3.3 Why Deep Neural Networks Did Not Break Through

The introduction of deep multi-layer perceptrons (MLPs) and recurrent sequence models appeared to dissolve the linearity constraint entirely. (Elhoseny *et al.*, 2022) reported 95.8% accuracy on financial distress prediction using a deeply optimised neural network — a figure that superficially suggests the prediction problem has been solved. It has not, for a reason that is worth spelling out carefully.

Financial distress prediction and M&A synergy prediction are fundamentally different questions. The former asks: “Is this firm, on its own, likely to fail?” That question can be answered, at least in part, from a single firm’s own financial time series. The latter asks: “Will combining these two specific firms create more value than they possess separately?” That is an inherently relational question — its answer depends on the interaction between the two firms and on their positions within a broader industrial ecosystem. An MLP with ten layers and ten thousand

parameters, trained on a vector of fifty financial ratios, still operates in a fifty-dimensional space that is topologically flat. It has no representation of whether the acquirer’s largest customer is about to go bankrupt, or whether the target and acquirer share a fragile single-source supplier. These omissions cannot be corrected by adding more layers to the same architecture.

Stream II conclusion: Every generation of tabular model — logistic regression, random forest, XGBoost, MLP, LSTM — inherits the independence assumption from its predecessor. The asymptotic ceiling is not a function of model complexity but of the information content of the feature space. Breaking through it requires a fundamentally different kind of model: one that operates on relationships rather than rows. Stream IV addresses this. Before arriving there, Stream III examines why adding language data — an intuitively promising additional information source — proved far more complicated than the field anticipated.

2.4 Stream III: The Semantic Turn — Powerful Tools, Wrong Targets

2.4.1 The Bag-of-Words Era: A Useful Foundation with Hard Limits

Where tabular models read financial numbers, semantic models read text — in the M&A context, primarily the annual 10-K reports that US-listed firms file with the Securities and Exchange Commission. These filings contain structured, standardised sections: the Management Discussion and Analysis (MD&A), in which management describes strategy, performance, and outlook in their own words, and the Risk Factors section, in which the company enumerates specific threats and uncertainties it faces. The intuition motivating semantic M&A research is that these disclosures might reveal complementarity or conflict between acquirer and target that financial ratios alone cannot capture.

The foundational contribution to financial NLP is (Loughran and McDonald, 2011), who demonstrated that standard general-purpose sentiment dictionaries systematically mislabel financial language. Words like “liability,” “risk,” and “obligation” carry negative connotations in ordinary English but are routine, neutral legal descriptors in financial filings. Their domain-specific word lists corrected this systematic mislabelling and established a principle that remains foundational: financial text requires tools calibrated for the specific conventions of financial language. This study directly inherits that principle in its choice of FinBERT over general-purpose language models.

The bag-of-words (BoW) paradigm that dominated early financial NLP, however, carries a structural limitation that no amount of vocabulary refinement can overcome: it is compositionally blind. A BoW model represents a document as a count of how many times each word appears, with no memory of the order or context in which words appeared. The sentence “We have eliminated our exposure to risk” and the sentence “We have significant exposure to risk” produce nearly identical word-frequency vectors while conveying directly opposite strategic signals. In M&A filings, where a single conditional clause — “subject to regulatory approval” — can invert the economic meaning of a multi-billion-dollar commitment, this ambiguity is not a minor inconvenience; it is a systematic failure mode.

Despite this, a substantial body of recent research has continued to extend the BoW approach: (Demers, Wang and Wu, 2024) construct frequency-based lexicons for human capital disclosures; (Acheampong *et al.*, 2025) proxy financial constraints using lexicon-based indices; (Garcia, Hu and Rohrer, 2020) develop finance-specific dictionaries validated through market data. Each of these contributions improves the vocabulary but cannot address the underlying compositional blindness. (Hoberg and Phillips, 2016) represents the most structurally sophisticated BoW application — their text-based network industries classification used product-description cosine similarity from 10-K filings to construct time-varying firm similarity networks, demonstrating that textual proximity predicts competitive dynamics and deal likelihood. This study inherits that semantic proximity methodology directly, but redirects it from predicting whether a deal will happen to predicting whether a deal that has happened will create or destroy value, and replaces TF-IDF word vectors with contextualised FinBERT embeddings.

2.4.2 The Transformer Revolution: A Better Tool, But Often Misdirected

(Devlin *et al.*, 2018) BERT architecture marked a genuine qualitative advance in language modelling. Rather than counting word occurrences, BERT learns contextual embeddings — vector representations in which the same word receives a different numerical encoding depending on the surrounding words. This directly resolves the compositional blindness of BoW models. (Araci, 2019) FinBERT specialised this architecture for financial language through pre-training on a large corpus of corporate filings, giving it domain-calibrated representations that a general BERT model lacks.

The challenge is not with the tools themselves but with how the field has applied them to M&A prediction. (Zhao, Li and Zheng, 2020) and (Lohmeier and Stitz, 2023) both used transformer architectures to extract semantic features from M&A-related text, but then compressed the resulting rich, high-dimensional representations into a single scalar sentiment score before feeding them into a downstream classifier. This pipeline — transformer as sophisticated word counter, followed by scalar compression — preserves almost none of the representational richness that motivated using a transformer in the first place. The architecture is, in terms of information content, not substantially different from an advanced bag-of-words approach.

More fundamentally, (Hájek and Henriques, 2024) and (Lohmeier and Stitz, 2023) both direct their NLP pipelines at predicting **deal occurrence** — whether an acquisition announcement will happen — rather than **deal outcome** — whether the announced deal will create or destroy shareholder value. These are genuinely different problems. Predicting deal occurrence requires detecting patterns in news flow that precede announcements; predicting deal outcome requires modelling the fundamental complementarity between acquirer and target across multiple dimensions. The existing NLP literature has made substantial progress on the former and has largely not addressed the latter.

2.4.3 Why Standard FinBERT Cannot Fully Recover the Textual Signal

Even correctly directed at deal outcome prediction, the standard single-document FinBERT implementation faces three structural limitations that this study directly addresses.

Frozen generic embeddings: FinBERT’s pre-training objective is general financial language modelling — predicting masked words in financial text — not the specific task of discriminating strategic alignment from risk concentration. Its frozen weights cannot adapt to the semantic distinctions most relevant to synergy prediction. This study addresses this by treating the two dominant 10-K sections — MD&A and Risk Factors — as semantically distinct modalities rather than interchangeable text, exploiting a distinction FinBERT cannot make on its own.

Single-document architecture: Studies that apply FinBERT independently to the acquirer’s filing or the target’s filing separately miss the cross-document signal that matters most for synergy prediction. It is not the absolute strategic sophistication of either party that predicts synergy — it is the **relative alignment** between the two, computed as the cosine distance between paired acquirer-target embeddings from the same document section. This pairwise computation is the core of H2 in this study.

Section conflation: Standard implementations embed entire 10-K documents or undifferentiated excerpts, mixing two sections that this study’s H2 predicts have **opposite** relationships with deal value. MD&A similarity (shared strategic direction) is hypothesised to correlate positively with announcement returns; Risk Factor similarity (overlapping liability exposure) is hypothesised to correlate negatively. A model that conflates these sections fits a noisy mixture of opposing signals, systematically suppressing the predictive coefficient of each. The -0.012 AUC drop observed when undifferentiated text embeddings are added to the model (M2 vs. M1) is the empirical consequence of exactly this conflation.

The choice of FinBERT over longer-context alternatives such as Longformer is justified on three grounds: domain-specific vocabulary alignment with 10-K filing language; the section-specific extraction protocol used here, which makes the 512-token context window a non-binding constraint in practice; and empirical evidence from (Qin and Yang, 2019) that frozen domain-specific representations with lightweight downstream prediction heads perform competitively on financial classification tasks at sample sizes comparable to this study’s deal universe.

Stream III conclusion: Transformer-based NLP has the representational capacity to capture strategic fit signals in 10-K filings. That capacity has not been applied to post-merger synergy outcome classification. Furthermore, even when correctly directed, text models operating on isolated documents miss the second-order structural reality: two firms with perfectly aligned strategies can still fail to create synergy if their network positions make the combined entity more fragile. This structural incompleteness is what Stream IV addresses.

2.5 Stream IV: The Topological Turn — Network Structure as Irreducible Signal

2.5.1 The Economic Foundation: Supply Chain Momentum

Before examining the empirical evidence, it is worth clarifying what “network topology” means in this context. Modern corporations do not exist in isolation — they are embedded in dense webs of economic relationships. A manufacturer buys components from dozens of suppliers; a

retailer sells through hundreds of partners; competitors share customers, regulatory exposure, and talent pools. These relationships create dependencies: if a key supplier fails, the customer's production line can stop; if a competitor is acquired and strengthens their market position, the remaining firm's pricing power erodes. Topology refers to the **structure** of these relationship networks — not just which firms are connected, but how centrally positioned each firm is, how many paths between other firms run through it, and whether it serves as a critical bridge between otherwise disconnected clusters. The central hypothesis of Stream IV is that this structural position encodes information about how a deal will be received by the market — information that balance sheets and management narratives cannot convey.

The empirical case for this hypothesis is grounded in a well-established finding from financial economics rather than in machine learning novelty. (Cohen and Frazzini, 2008) demonstrated that economic shocks to a supplier do not immediately price into the stock of the supplier's customer. Due to information friction and limited investor attention, these shocks propagate across supply chain links with a measurable time lag, generating a predictable return momentum effect for investors who monitor inter-firm dependency structures. This “supply chain momentum” phenomenon is direct empirical evidence that network topology encodes information about future firm value that is absent from standalone financial metrics — and that this information persists long enough to be exploited, precisely because most models process firms in isolation. (Ahern and Harford, 2014) extended this reasoning directly to M&A. They demonstrated that the structure of industry-level trade networks predicts the direction of merger waves: acquisitions cluster along supply chain linkages because acquiring firms seek to internalise relationships that generate the highest dependency-reduction value. More directly relevant to this study, they found that supply chain proximity at the industry level predicts post-merger combined stock returns — providing precedent that network position carries synergy-relevant information over and above financial fundamentals. This study operationalises the same insight at the firm level using Bloomberg SPLC data, enabling node-level rather than industry-level topology encoding.

2.5.2 Corporate Network Centrality and Value

The broader corporate networks literature reinforces this foundation. (Fee and Thomas, 2004) documented that customer-supplier relationships carry systematic pricing power information: firms with high customer concentration earn measurably different returns upon resolution of those relationships, confirming that vertical network dependencies have real equity value implications. (Larcker, So and Wang, 2013) demonstrated that director network centrality — specifically betweenness centrality in board interlock networks — is a significant predictor of future firm performance and acquisition premia, establishing precedent that graph-theoretic centrality measures carry economic signal in corporate finance contexts.

To understand betweenness centrality intuitively: imagine a social network where some people are highly connected to many different groups while others only know people within their own immediate circle. Betweenness centrality measures how often a given node sits on the shortest path between two other nodes — in other words, how much of the network's “traffic” must

pass through it. A firm with high betweenness centrality in the supply-chain-competitor graph is a structural broker: many supplier-customer or competitor-competitor connections in the industrial ecosystem run through it. These “bridge nodes” face a distinctive dual constraint: their centrality creates both upside synergy capture opportunities (they can internalise strategically important relationships) and downside disruption exposure (shocks from any connected cluster propagate to them quickly). This bilateral constraint is the mechanism hypothesised in H3 to compress the variance of announcement-period returns relative to peripheral acquirers.

2.5.3 Why Adding Centrality to XGBoost Is Insufficient

A sophisticated critic might ask: if betweenness centrality is valuable information, why not simply compute it and add it as an additional column in the feature table? This scalar centralisation approach has two structural limitations that motivate a full graph neural network architecture.

First, scalar centrality measures collapse a high-dimensional structural signal into a single number, discarding the specific **pattern** of connectivity that carries economic meaning. Two firms can share identical betweenness centrality scores while occupying radically different network contexts: one may bridge two healthy, high-growth industrial clusters; the other may bridge two sectors in secular decline. A scalar feature treats these as identical; a graph embedding that propagates and aggregates neighbourhood information preserves the distinction (Hamilton, Ying and Leskovec, 2017).

Second, scalar features cannot represent second-order network effects — dependencies that run through intermediate nodes. If an acquirer’s primary supplier is itself a customer of the target, this creates a circular dependency loop that substantially changes the risk profile of the combined entity. This loop is invisible to any feature engineering pipeline that processes each deal in isolation. A message-passing graph neural network propagates information through multi-hop neighbourhoods by design, recovering these higher-order structural patterns without requiring their explicit enumeration (Hamilton, Ying and Leskovec, 2017).

(Venuti, 2021) provided practical validation of this argument: GraphSAGE achieved 81.79% accuracy in predicting acquisitions for private companies where financial data is sparse — precisely the regime where structural neighbourhood information provides the largest marginal value over standalone features. This establishes the inductive architecture’s applicability to data-scarce M&A contexts.

2.5.4 Why Edge-Type Semantics Matter: The Case for Heterogeneous Graphs

A final refinement is necessary before the graph architecture can be considered adequate. A homogeneous graph model — one that treats all inter-firm edges as equivalent — commits a category error that compromises its predictions.

A `supplies_to` relationship implies operational dependency and risk propagation: if the supplier fails, the customer’s production line stops. A `competes_with` relationship implies market concentration and pricing power dynamics: if two competitors merge, their combined

market share may attract regulatory intervention and change the competitive landscape for all remaining players. These are fundamentally different economic mechanisms. Collapsing them into a single undirected edge type introduces a mixed signal that obscures what each relationship type independently encodes — analogous to averaging the coefficients of variables that have opposite effects.

(Shi *et al.*, 2017) formalised the theoretical basis for Heterogeneous Information Networks (HINs) and demonstrated that type-specific attention mechanisms during neighbourhood aggregation consistently outperform homogeneous baselines on multi-relational graphs. (Wang *et al.*, 2021) extended this into the Heterogeneous Graph Attention Network (HAN), showing that relationship-type-specific aggregation recovers semantic structure that single-type models discard. (Lv *et al.*, 2021) provided systematic benchmarks confirming heterogeneous graph encoders achieve statistically superior performance across multiple classification tasks on multi-relational corporate networks.

While theoretical heterogeneous graph frameworks $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathcal{T}_v, \mathcal{T}_e)$ can encode four distinct edge relationships (supplier_of, customer_of, competitor_of, and acquires), this study explicitly restricts the implemented architecture to only two active supply-chain dependencies (supplies and buys_from). As detailed in Chapter 3, this restriction is structurally necessary to prevent target leakage from competitor or acquisition edges. Type-specific GraphSAGE aggregation functions are applied independently per relationship type before cross-type pooling — recovering exactly the semantic distinctions that scalar feature engineering and homogeneous graph models discard.

Stream IV conclusion: Graph neural network architectures are theoretically well-suited to encoding inter-firm dependencies, and supply chain network structure is empirically demonstrated to contain synergy-relevant information. Yet no published study has directed a heterogeneous GNN at post-merger synergy **outcome** classification using firm-level topology. The Topological Alpha Hypothesis (H1), the Topological Arbitrage Hypothesis (H3), and the HeteroGraphSAGE architecture proposed in this study constitute the original contribution at this intersection.

2.6 The Measurement Foundation: Event Study Methodology

2.6.1 How Announcement Returns Are Measured

While Streams I through IV establish the feature categories required to **predict** synergy, a rigorous framework is required to **measure** synergy itself — a non-trivial problem given that the true economic benefits of a merger may take years to fully materialise. The standard solution in empirical corporate finance, formalised by (MacKinlay, 1997) and whose statistical properties were rigorously characterised by (Brown and Warner, 1985), is the event study.

The core idea is elegant: if financial markets are reasonably efficient at incorporating publicly available information, then the stock price reaction to a deal announcement reflects the market’s best estimate of the present value of future synergies. A large positive price reaction means the

market believes the deal will create substantial value; a negative reaction means it will destroy value.

The methodology operationalises this through a market model. For each acquirer i , a “normal” return generating process is estimated over a pre-announcement estimation window:

$$R_{i,t} = \alpha_i + \beta_i R_{m,t} + \varepsilon_{i,t}$$

where $R_{i,t}$ is the firm’s daily return, $R_{m,t}$ is the broad market return, and α_i and β_i are firm-specific parameters. Abnormal Returns (ARs) during the announcement period are then computed as the difference between what the firm actually earned and what the model predicted it should have earned given market conditions:

$$AR_{i,t} = R_{i,t} - (\hat{\alpha}_i + \hat{\beta}_i R_{m,t})$$

The Cumulative Abnormal Return (“CAR”) sums these daily excess returns across the event window:

$$CAR_i = \sum_{t=t_1}^{t_2} AR_{i,t}$$

A positive CAR indicates the market interpreted the deal as value-creating; a negative CAR indicates value destruction. This is the binary outcome label — positive or negative — that all models in this study attempt to predict.

2.6.2 Event Window Choice

(MacKinlay, 1997) established that event window selection involves a fundamental bias-variance trade-off. Narrow windows (e.g., $[-1, +1]$ trading days) isolate the immediate announcement surprise cleanly but risk missing delayed market reactions for deals where information leaked before the formal announcement. Wider windows (e.g., $[-5, +5]$) capture more complete price discovery but introduce additional variance from unrelated news.

This study uses the $[-5, +5]$ window, consistent with the empirical M&A literature (Betton, Eckbo and Thorburn, 2008), as a practical balance that accommodates the pre-announcement information leakage documented to begin three to five days before announcement in approximately 25% of transactions. While a narrower $[-1, +1]$ specification could theoretically isolate the immediate announcement shock more cleanly, the wider window is necessary here to ensure that delayed price discovery in smaller, less liquid acquirers is fully captured.

2.6.3 Why Binary Classification Rather Than Regression

The signal-to-noise characteristics of short-window CARs make predicting their **magnitude** a challenging target for machine learning models. Announcement return magnitudes are driven by surprise — and surprise, by definition, is hard to predict from publicly available information. The consistently low pseudo- R^2 values documented in the literature (Palepu, 1986; Betton,

Eckbo and Thorburn, 2008) are consistent with the Efficient Markets Hypothesis’s implication that the **size** of the market’s reaction to new information is largely unpredictable.

The **direction** of the reaction, however, contains learnable signal rooted in deal fundamentals: whether the acquirer has the financial capacity to absorb the target, whether their strategies are complementary, and whether their ecosystem positions make the combination structurally sound. This study therefore defines $y_i = 1$ if $CAR_i > 0$ (value-creating) and $y_i = 0$ otherwise, with AUC-ROC as the primary evaluation metric for its robustness to class imbalance and threshold-invariant characterisation of discriminative power (Ajayi *et al.*, 2022).

2.7 The Multimodal Imperative: Why Fusion Is the Only Exit

2.7.1 Complementary Variance: Three Lenses on the Same Deal

The four streams surveyed above converge on a structural finding: financial, textual, and topological features encode different, partially non-overlapping dimensions of synergy potential. (Baltrušaitis, Ahuja and Morency, 2019) formalize this as “complementary variance” in their taxonomy of multimodal learning architectures. A firm may exhibit strong capital adequacy (visible in financial ratios), precise strategic alignment with its target (visible in text), yet suffer from fragile single-source supplier dependencies (visible only in graph topology). A model that observes only one of these dimensions is mathematically blind to whatever divergences exist in the others, and no within-modality engineering can recover information that was never measured.

(Xu *et al.*, 2021) demonstrated the practical validity of this in a financial forecasting context: structured financial data and unstructured text features carry statistically independent predictive signal — signal that vanishes from each modality when the other is controlled for, but is recoverable when both are jointly modelled. (Qin and Yang, 2019) demonstrated the same complementarity between price-based and news-based features for stock movement prediction.

2.7.2 Why HeteroGraphSAGE Fusion Is the Logical Response

Standard baselines — logistic regression on financial ratios, XGBoost on financial ratio vectors, standalone FinBERT applied to individual documents, homogeneous GCN with scalar centrality features — each fail for a specific, architecturally irreversible reason. Table 2 maps each study to its structural failure mode and the component of this study’s architecture designed to address it.

Literature Stream	Legacy Method	Structural (The Ceiling)	Failure	Implemented Architecture (The Exit)
Stream I: The M&A Paradox & Valuation				
Traditional Corporate Finance (Roll, 1986; Sirower, 1997)	DCF / Valuation Models	Behavioural contamination; hubris; market-timing noise	contaminated executive market-timing	Cumulative Abnormal Return (CAR) as objective market target
Stream II: The Tabular Paradigm				
Palepu (1986); Barnes (1990)	Logit / MDA on ratios	Linearity bias; pseudo- $R^2 < 0.10$; no network features		Non-linear fusion engine; Blocks A+B+C
Zhang et al. (2024)	XGBoost / Random Forest	i.i.d. assumption; survivorship bias; look-ahead contamination		GraphSAGE neighbourhood aggregation; temporal splits
Elhoseny et al. (2022)	Deep MLP (AWOA-DL)	Topologically flat; predicts financial distress, not synergy		Relational graph operator; synergy CAR as target
Stream III: The Semantic Turn				
Loughran & McDonald (2011)	Bag-of-Words (BoW)	Compositional blindness; deal occurrence target		Contextual embeddings; section-specific similarity (H2)
Hajek et al. (2024)	FinBERT sentiment	Single-document; no pairwise delta; acquisition likelihood target		Pairwise acquirer-target cosine distance; H2 conditional directionality
Zhao et al. (2020); Han et al. (2023)	BERT / RoBERTa → XGBoost	Conflates strategic and risk disclosures; predicts occurrence not outcome		Section-specific semantic splitting; binary CAR classification (H2)

Table 1: Methodological Evolution Matrix (Part I): Evolution of M&A prediction research through three paradigms.

Literature Stream	Legacy Method	Structural (The Ceiling)	Failure	Implemented Architecture (The Exit)	Archi-
Stream IV: Topology and Multimodal Fusion					
Cohen & Frazzini (2008); Ahern & Harford (2014)	Industry-level supply chain analysis	Industry-level only; no firm-level GNN; no CAR prediction		Firm-level SAGE on SPLC data (H1, H3)	Graph-
Venuti (2021)	Homogeneous GraphSAGE	Collapses edge semantics (suppliers = competitors); deal likelihood target		Heterogeneous edge types; full multimodal fusion	
Baltrušaitis et al. (2019); Xu et al. (2021)	Multimodal fusion frameworks	Not applied to M&A synergy; no modality		Block A+B+C late fusion with SHAP decomposition	

Table 2: Methodological Evolution Matrix (Part II): The transition to topological and multimodal fusion.

The HeteroGraphSAGE fusion architecture resolves each limitation simultaneously: GraphSAGE’s inductive neighbourhood aggregation recovers the topological signal that tabular models cannot access; FinBERT’s contextual embeddings on section-split 10-K filings recover the semantic signal that BoW and scalar NLP cannot represent; late fusion via concatenation preserves modality-specific representations while enabling cross-modal learning in the joint prediction head; and heterogeneous edge-type encoding preserves the semantic distinctions between relationship types that homogeneous graph models discard.

2.7.3 Architectural Choice: Late Fusion and Its Rationale

The M&A deal universe with complete multimodal coverage is approximately 2,800–5,000 observations — substantially below the sample sizes typically required for stable joint end-to-end training of transformer and GNN components (Baltrušaitis, Ahuja and Morency, 2019). This study therefore adopts late fusion: each modality is encoded independently into a fixed-dimensional embedding vector ($\mathbf{h}_F \in \mathbb{R}^{d_F}$, $\mathbf{h}_T \in \mathbb{R}^{d_T}$, $\mathbf{h}_G \in \mathbb{R}^{d_G}$), and these vectors are concatenated into a joint representation $\mathbf{z}_i = [\mathbf{h}_F \parallel \mathbf{h}_T \parallel \mathbf{h}_G]$ before a shared prediction head. This design isolates modality-specific representation learning from cross-modal inference, enabling robust individual stream training even when certain modalities have incomplete coverage.

The downstream prediction head uses (Chen and Guestrin, 2016) XGBoost framework for the baseline ablation experiments, given its well-documented advantages over neural networks on heterogeneous tabular feature vectors at M&A sample sizes. SHAP decomposition (Lundberg and Lee, 2017) is applied post-inference to provide per-modality attribution scores, enabling the ablation experiments to quantify the marginal contribution of each block to predictive accuracy.

2.8 Synthesis: Gap Table and Research Hypotheses

No prior published study has jointly fused financial fundamentals, FinBERT textual embeddings, and supply-chain graph topology, directed this multimodal architecture at binary CAR direction classification as a measure of post-merger synergy, and tested the marginal contribution of each modality through controlled ablation. The gap table above maps the precise points at which prior work reaches its limits. The three formal hypotheses below emerge directly from those limits.

2.8.1 Formal Research Hypotheses

H1 — The Topological Alpha Hypothesis: The inclusion of second-order neighbour embeddings via GraphSAGE (Block C) will yield a statistically significant increase in AUC-ROC relative to the financial-only baseline (Block A), under five-fold stratified cross-validation. This gain will be disproportionately concentrated within supply-chain-dependent manufacturing sectors (SIC 20–49) compared to asset-light service sectors (SIC 60–79), reflecting the hypothesis that graph embeddings recover signal proportional to industrial ecosystem structural density (Cohen and Frazzini, 2008; Ahern and Harford, 2014).

H2 — The Semantic Divergence Hypothesis: The predictive relationship between textual similarity and synergy direction is conditional on document section.

$$\text{CAR}_i = \beta_0 + \beta_1 \cdot \text{sim}_{\text{MDA},i} + \beta_2 \cdot \text{sim}_{\text{RF},i} + \varepsilon_i$$

and directly tests the monotonic sentiment utility assumption embedded in standard NLP pipelines (Loughran and McDonald, 2011; Hájek and Henriques, 2024).

H3 — The Topological Arbitrage Hypothesis: Acquirer nodes with high betweenness centrality in the heterogeneous supply-chain graph will exhibit statistically compressed variance in $|\text{CAR}|$ outcomes relative to peripheral nodes, as measured by Levene’s test across betweenness centrality quantile groups. This tests the Information Transparency Dampening mechanism: structurally prominent acquirers are more continuously monitored by the market, making their deal announcements less surprising and their returns less dispersed (Larcker, So and Wang, 2013; Ahern and Harford, 2014).

Chapter 3: Methodology

3.1 Introduction

This chapter details the research design, architectural implementation, and evaluation protocols employed to construct the proposed tri-modal fusion model for M&A synergy prediction. The methodology is structured to operationalise the theoretical findings from Chapter 2, translating the need for “multimodal fusion” into a rigorous engineering specification.

The chapter proceeds as follows. Section 3.2 establishes the research philosophy and epistemological stance. Section 3.3 formalises the three hierarchical hypotheses. Section 3.4 describes all data sources and collection pipelines. Section 3.5 covers cleaning, normalisation, and temporal splitting. Section 3.6 defines the three feature blocks (Financial, Textual, Graph). Section 3.7 specifies the baseline and fusion model architectures. Section 3.8 derives the CAR target variable. Section 3.9 details the hypothesis-testing protocol, and Section 3.11 addresses limitations and ethical considerations.

3.2 Research Philosophy and Design




Figure 1: System architecture pipeline: Multimodal feature extraction, fusion, and evaluation.

This study adopts a **post-positivist** epistemological stance, treating M&A synergy as a latent, probabilistic construct approximated through market reactions and structured inter-firm relationships. While acknowledging that markets are not perfectly efficient, the research operates within the semi-strong form of the Efficient Market Hypothesis (Fama, 1970; 1991), wherein publicly available information — financial fundamentals, regulatory filings, and network topology — constitutes a viable predictor signal (Jensen, 1978). The overarching research design is **quantitative and deductive**: three *a priori* hypotheses (H1, H2, H3) are specified before analysis and tested through controlled ablation experiments (Creswell, 2014).

The study employs a **cross-sectional observational design**. Because M&A deals are historical and non-repeatable, no experimental manipulation is possible; causal inference is instead approximated through systematic covariate control, ablation modelling, and statistical hypothesis testing, following standard practice in empirical corporate finance (MacKinlay, 1997).

An **Experimental Prototyping SDLC** governs the engineering programme. Non-deterministic model outputs and stochastic training dynamics demand reproducible seeding, versioned artefacts, and isolated ablation configurations rather than a traditional waterfall build process. Each experimental variant is fully specified in a YAML configuration (e.g., `full_fusion.yaml`) that pins hyperparameters, random seeds, feature subsets, and evaluation splits.

3.3 Research Hypotheses

Three hierarchical hypotheses structure the empirical programme. Together they form a logical escalation: H1 tests whether the graph stream improves directional discrimination (classification) above financials; H2 tests whether pairwise textual similarity carries directionally opposing signal depending on filing section; H3 tests whether graph centrality compresses the variance of announcement returns via Levene’s test across centrality quantile groups.

Predicting both the sign and the magnitude of CAR serves two distinct analytical purposes. Predicting the **magnitude** (via regression) is notoriously difficult due to competing bids, information asymmetry, and macro shocks, but it is necessary to test structural variance compression (H3). Predicting the **sign** (via classification) is highly tractable and systematically links to practical deal advisory: whether the deal is a net positive or a net negative (H1, H2). The architecture therefore formulates CAR as a **dual target**: a continuous variable for variance analysis, and a thresholded binary label (1 if $CAR > 0$, 0 otherwise) for directional discrimination.

ID	Name	Formal Statement
H1	Topological Alpha	Supply-chain network centrality metrics derived from Bloomberg SPLC carry statistically significant predictive signal for acquirer CAR direction, incremental to financial fundamentals alone (statistically significant AUC-ROC improvement on held-out test set, $p < 0.05$ by paired t -test across cross-validation folds).
H2	Semantic Divergence	The cosine distance between acquirer and target FinBERT embeddings of their respective 10-K MD&A sections is a significant predictor of post-acquisition CAR ($p < 0.05$, Pearson/Spearman correlation; confirmed by ablation).
H3	Topological Arbitrage	Acquirer nodes with high betweenness centrality in the heterogeneous supply-chain graph will exhibit statistically compressed variance in $ CAR $ outcomes relative to peripheral nodes. Success criterion: Levene’s test ((Levene, 1960)) for equality of variances across centrality quantile groups yields $p < 0.05$, confirming the Information Transparency Dampening mechanism.

Table 3: Research Hypotheses

3.4 Data Sources and Collection

3.4.1 M&A Deal Universe

The primary dataset is sourced from **Yahoo Finance**, which provides deal-level financial attributes for completed M&A transactions. completed M&A transactions. Five raw CSV exports are merged via `scripts/data/build_combined_dataset.py` into `data/interim/ma_combined.csv`.

The deal universe is restricted to:

- Completed acquisitions of publicly listed US targets by publicly listed US acquirers.
- Transactions announced between 2000 and 2023.
- Deal values exceeding USD 50 million (sufficient market microstructure data for reliable CAR estimation).

These filters follow established practice (Betton, Eckbo and Thorburn, 2008) and ensure a minimum of 120 trading days in the estimation window.

3.4.2 Equity Return Data

Daily equity returns for acquirer firms and the S&P 500 benchmark are retrieved via **yfinance** (Bloomberg ticker conversion handled by `pull_car_data.py`). Returns are aligned to deal announcement dates and stored in a long-format time series (`timeseries_long.csv`) with a `rel_day` field denoting trading-day distance from announcement (Day 0 = first trading day on or after the announcement date, forward-fill rule). Failed ticker lookups are retried with fuzzy-matching heuristics.

3.4.3 Textual Data (SEC EDGAR)

10-K annual filings for acquirer firms are retrieved from the **SEC EDGAR full-text search API**, targeting the MD&A (Item 7, `item_7_mda.txt`) and Risk Factors (Item 1A, `item_1a_risk.txt`) sections for the fiscal year immediately preceding each announcement. Extraction is handled by `src/features/edgar.py`, with download provenance logged in `data/external/edgar/download_log.csv`.

3.4.4 Supply-Chain Network Data

Inter-firm supply-chain relationships are sourced from **Bloomberg SPLC** (Supply Chain Analysis), which maps disclosed customer–supplier relationships for publicly listed firms. The SPLC data is merged with the deal universe via `scripts/data/merge_splc_data.py`, matching on Bloomberg ticker symbols. This forms the edge set for the heterogeneous graph constructed in Section 3.6.3.

3.5 Data Preprocessing

3.5.1 Cleaning and Quality Control

Raw Yahoo Finance exports undergo systematic cleaning in `scripts/data/data_cleaning.py`: date parsing and standardisation; deduplication of records sharing the same acquirer–target–announcement-date triplet; removal of records with missing acquirer ticker or announcement date; and currency normalisation to USD using period-end exchange rates.

3.5.2 Feature Engineering and Normalisation

Financial features comprise 56 ratio-level variables spanning acquirer and target leverage, liquidity, profitability, and deal structure characteristics. The preprocessing pipeline applies:

1. **Winsorisation** at the 1st and 99th percentile to bound outlier influence.
2. **Z-score standardisation** (zero mean, unit variance) computed on training-set statistics *only*, then applied to validation and test sets.
3. **Stratified temporal splitting** (70 / 15 / 15) by announcement year to prevent temporal leakage.

3.5.3 Temporal Splitting and Event-Window Embargo

Deals are sorted chronologically and partitioned into training (2000–2016), validation (2017–2019), and test (2020–2023) sets based on announcement year. This strict temporal ordering ensures the model never trains on information post-dating the validation or test periods.

Temporal Dataset Partition

Partition	Years	Split	Role
Training	2000–2016	70%	Model fitting
Validation	2017–2019	15%	Hyperparameter tuning + early stopping
Test	2020–2023	15%	Final held-out evaluation

Embargo: An 11-trading-day gap is enforced at each temporal boundary. Any deal announced within 11 days of a split boundary is excluded from both adjacent partitions, preventing CAR event-window overlap (cf. (López de Prado, 2018)).

Table 4: Temporal splits and embargo design

Concretely, because the event window spans $[-5, +5]$ trading days, two deals whose announcement dates differ by fewer than 11 trading days share overlapping market-return sequences in their CAR calculations. If one such deal falls in the training set and the other in validation, the model can implicitly learn return correlations that exist only because of calendar proximity — the **Overlapping Outcomes** problem formalised by (López de Prado, 2018). The 11-day embargo eliminates this cross-contamination by construction.

3.5.4 Missing Data Strategy

Features with $> 40\%$ missing values are excluded. For remaining missing values, median imputation is applied to continuous features and mode imputation to categorical indicators. All imputation statistics are fitted on the training set only.

3.6 Feature Extraction

The three feature blocks are summarised in Table 5.

Block	Modality	Source	Construction (Fusion Pipeline)
A	Financial	Yahoo Finance	56-column ratio matrix → Winsorise → z-score → ProjectionHead ($\mathbb{R}^{56} \rightarrow \mathbb{R}^{64}$, linear + ReLU)
B	Textual (FinBERT)	SEC EDGAR 10-K (Item 7 + Item 1A)	FinBERT tokenisation (512-token chunks, stride = 256) → [CLS] from penultimate layer → mean-pool across chunks → $\mathbb{R}^{768}$ per section → PCA compression (fit on train only): $\mathbb{R}^{768} \rightarrow \mathbb{R}^{64}$ per section → concatenate MD&A + RF vectors → $\mathbb{R}^{128}$ total. <i>Note: pairwise cosine similarity between acquirer and target section embeddings is computed separately as a scalar predictor for the H2 OLS test; it is not an input to the fusion model.</i>
C	Graph (HeteroGraph-SAGE)	Bloomberg SPLC	HeteroConv (2-layer SAGE-Conv, separate per edge type) on (company, supplies, company) and (company, buys_from, company) edges → 64-dim node embedding per firm. <i>Deals without SPLC coverage receive a zero vector; the graph stream is masked via has_graph = False in the fusion model.</i>

Table 5: Feature block definitions and fusion-pipeline construction steps.

3.6.1 Block A — Financial Features

The financial feature vector $\mathbf{h}_F \in \mathbb{R}^{d_F}$ is constructed directly from the standardised preprocessing output ($d_F = 56$). For baseline models (Ridge Regression, ElasticNet, XGBoost), $\mathbf{h}_F$ is used directly. For the MLP and fusion models, it passes through a ProjectionHead — a linear layer followed by ReLU — that maps $\mathbf{h}_F$ to a lower-dimensional embedding $\hat{\mathbf{h}}_F \in \mathbb{R}^{64}$ before concatenation.

3.6.2 Block B — Textual Features (FinBERT)




Figure 2: Section-aware FinBERT NLP pipeline extracting MD&A and Risk Factors semantic vectors.

Each acquirer firm’s MD&A (Item 7) and Risk Factors (Item 1A) text is processed through **FinBERT** (ProsusAI/finbert) (Araci, 2019), a BERT-base architecture fine-tuned on financial communications corpora. The exact pipeline, as implemented in `src/features/text.py`, is:

1. **Chunking.** The raw section text is tokenised (no truncation) and split into overlapping 512-token windows with a stride of 256 tokens, reserving positions 0 and 511 for the [CLS] and [SEP] tokens respectively.
2. **Extraction.** For each chunk, the [CLS] token representation is taken from the **penultimate transformer layer** (`hidden_states[-2]`), yielding a 768-dimensional vector.
3. **Pooling.** Chunk-level vectors are **mean-pooled** across all chunks to produce a single $\mathbf{h}_T \in \mathbb{R}^{768}$ per section.
4. **PCA compression.** Each section’s embedding matrix is independently PCA-compressed:

$$\mathbf{h}_T^{\text{section}} \in \mathbb{R}^{768} \xrightarrow{\text{PCA, fit on train only}} \mathbf{p}^{\text{section}} \in \mathbb{R}^{64}$$

Separate PCA models are fitted for MD&A and Risk Factors, serialised to `data/processed/pca_models.pkl` for reproducible inference. This compression reduces the 1536-dimensional raw concatenation to 128 dimensions while retaining maximum explained variance.

1. **Concatenation.** The two 64-dimensional section vectors are concatenated into the final textual embedding:

$$\mathbf{h}_T = [\mathbf{p}^{\text{MDA}} \parallel \mathbf{p}^{\text{RF}}] \in \mathbb{R}^{128}$$

FinBERT's $\approx 110\text{M}$ parameters are **frozen** throughout all downstream training to prevent overfitting given the limited M&A sample size; only the downstream projection heads are trained.

3.6.2.1 Cosine Similarity (H2 Test Only)

For the H2 semantic-divergence hypothesis, a pairwise cosine similarity score is computed **separately** between the acquirer and target's section embeddings, *after* PCA compression:

$$\text{SemanticDiv}_i = 1 - \frac{\mathbf{p}_{\text{acq}}^{\text{MDA}} \cdot \mathbf{p}_{\text{tgt}}^{\text{MDA}}}{\|\mathbf{p}_{\text{acq}}^{\text{MDA}}\| \cdot \|\mathbf{p}_{\text{tgt}}^{\text{MDA}}\|} \quad (\text{Salton and McGill, 1983})$$

This scalar divergence score is used **exclusively** as the independent variable in the H2 OLS regression. It is **not** an input to the fusion model. The distinction is critical: cosine distance is a *relationship-level* scalar characterising strategic fit between two firms, while the fusion model requires *firm-level* vectors to learn independent acquirer representations.

3.6.3 Block C — Graph Features (HeteroGraphSAGE)




Figure 3: HeteroGraphSAGE embedding process for supply-chain topological representation.

The inter-firm supply-chain network is constructed as a heterogeneous graph $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathcal{T}_v, \mathcal{T}_e)$ from Bloomberg SPLC data, using PyTorch Geometric's HeteroData object (`scripts/graphs/build_hetero_graph.py`).

3.6.3.1 Graph Structure




Figure 4: Supply-chain graph construction showing inter-firm dependencies and topological topology.

- **Node type.** A single node type company represents each publicly listed firm present in the SPLC dataset. Node features are initialised with degree centrality, betweenness centrality, and the standardised financial feature vector.
- **Edge types.** Two directed relationship types are encoded:

Edge Type	Direction	Semantics
supplies	supplier → customer	Firm A discloses Firm B as a customer in SPLC; directional dependency.
buys_from	customer → supplier	Inverse of supplies encodes up-stream procurement risk.

All edges represent **currently active, disclosed supply-chain relationships** sourced from SPLC. No M&A-derived edges (e.g., historical acquisition links) are included in the adjacency matrix. This design choice eliminates any possibility of structural target leakage: the model has no graph path connecting an acquirer to its deal target, because no such edge class exists.

Leakage Note. The absence of acquisition edges is intentional and architecturally guarantees zero structural target leakage. Supply-chain relationships exist independently of M&A outcomes and persist whether or not a deal completes. The model cannot “see” the deal being predicted through the graph.

Four edge types were considered in the project design phase (supplier_of, customer_of, competitor_of, acquires). The implemented scope uses two SPLC-sourced types. Competitor and historical acquisition edges are a natural extension discussed in Section 3.11.

3.6.3.2 HeteroGraphSAGE Model

A 2-layer **Heterogeneous GraphSAGE** model (Hamilton, Ying and Leskovec, 2017) is trained via self-supervised link prediction on the supply-chain graph, as implemented in `scripts/graphs/train_hetero_graph.py`:

HeteroGraphSAGE Architecture		
Layer	Operation	Detail
Conv 1	HeteroConv	Separate SAGEConv(in $\rightarrow$ 128) per edge type; mean aggregation across types
Activation	ReLU + Dropout	$p = 0.3$; applied per-type after Layer 1
Conv 2	HeteroConv	Separate SAGEConv(128 $\rightarrow$ 64) per edge type; mean aggregation
Output	Node embedding	$\mathbf{h}_G \in \mathbb{R}^{64}$ per company node

Training: self-supervised link prediction (binary cross-entropy), negative sampling per edge type, Adam lr = 0.01, 200 epochs.
Final embeddings extracted via `model.encode()` with full edge set.

Table 6: HeteroGraphSAGE architecture and training configuration.

The key architectural innovation over a homogeneous GraphSAGE is that `supplies` and `buys_from` edges learn **independent SAGEConv weight matrices**, enabling the model to distinguish upstream procurement signals from downstream customer dependency signals during message passing. Node embeddings are extracted and stored in `data/interim/`

hetero_graph_embeddings.csv (64-dimensional, one row per company ticker), then merged into the training dataset via deal–ticker matching.

3.7 Model Architecture

3.7.1 Baseline Models

Four baselines are trained on Block A features only:

Model	Variant ID	Purpose
Financial Only	M1	Non-linear baseline (XGBoost); tests H1/H2 incremental gain.
Financial + Text	M2	Ablation: tests H2 in isolation; verifies text aggregation effect.
Full Fusion (F+T+G)	M3	Primary model; tests H1.
Full Fusion + Aux	M3e	Extended model with auxiliary engineered features. ¹

Table 7: Model variants for ablation experiments.

¹M3e extends M3 by appending 13 scalar auxiliary features to the concatenated embedding vector z_i : two deal-level derived features (relative deal size; cross-sector indicator), six graph-theoretic centrality scalars for acquirer and target nodes (betweenness centrality, degree centrality, PageRank), three pairwise textual similarity scalars (MD&A cosine similarity, Risk Factor cosine similarity, cross-section divergence), one acquirer-to-target asset size ratio, and the payment method indicator. These scalars are computed prior to the ProjectionHead and appended directly to $z_i \in \mathbb{R}^{261}$. M3e does not improve AUC beyond M3 (0.5585 vs 0.5655), confirming that scalar centrality and similarity features add classification noise rather than signal when graph neighbourhood embeddings are already present.

3.7.2 Fusion Model Architecture




Figure 5: Multimodal late-fusion architecture combining financial, textual, and graph embeddings.

The primary model is the **late-fusion tri-modal architecture** implemented in `src/models/fusion.py`. Each active stream passes through its own ProjectionHead (linear + ReLU), and the resulting embeddings are concatenated:


$$z_i = [h_F \parallel h_T \parallel h_G] \in \mathbb{R}^{d_{F'} + d_{T'} + d_{G'}}$$

where $d_{F'} = 64$, $d_{T'} = 64$, $d_{G'} = 32$ by default. To rigorously evaluate both the magnitude of synergy and the practical investment decision, the architecture implements a **Dual-Evaluation Framework**. The concatenated feature vector z_i serves as the input to two parallel, independent training pipelines:

- **Regressor Pipeline (H3):** Includes a two-layer MLP with a **linear** output activation (and XGBRegressor), optimised via Mean Squared Error (MSE) to predict continuous CAR. This tests structural variance explanations.

- **Classifier Pipeline (H1, H2):** Includes a two-layer MLP with a **sigmoid** output activation (and XGBClassifier), optimised via Binary Cross-Entropy (BCE) to predict the binary direction of CAR (1 if CAR > 0, 0 otherwise). This tests practical deal advisory discrimination.

Both pipelines share the same upstream concatenated representation z_i , meaning the feature streams are extracted once and evaluated across both objectives.



Listing 1: Tri-modal late-fusion architecture (src/models/fusion.py).


3.7.3 Training Configuration

All PyTorch models are trained with:

- **Optimiser:** AdamW with cosine annealing learning-rate schedule with warm restarts.
- **Loss:** Two objectives trained in parallel: MSE for the Regressor Pipeline (continuous CAR); Binary Cross-Entropy for the Classifier Pipeline (binary CAR direction). Huber loss sensitivity analysis is reported alongside MSE.
- **Early stopping:** Validation MAE with patience of 15 epochs.
- **Batch size:** 64.
- **Reproducibility:** Fixed random seed via `set_seed()` in `src/training/trainer.py`.
- **Device:** CUDA / Apple MPS / CPU auto-selected via `src/config.py`.

3.8 Target Variable: Cumulative Abnormal Return

The target variable y_i for each deal is the **Cumulative Abnormal Return (CAR)** over the symmetric event window $[-5, +5]$ trading days relative to announcement date, computed via the standard market model (Fama and MacBeth, 1973; Brown and Warner, 1985; MacKinlay, 1997). This section describes the full two-stage pipeline: Stage 1 derives **actual** CAR values from raw market data using OLS; Stage 2 trains the fusion model to **predict** CAR from pre-announcement features and evaluates predicted CAR against actual CAR on held-out deals.



Listing 2: The two-stage CAR pipeline: Stage 1 computes actual CAR via OLS event study; Stage 2 trains and evaluates the fusion model against those labels.

3.8.1 Stage 1: OLS Market Model

The market model is estimated over the **estimation window** $[-200, -20]$ trading days (180-day window, minimum 120 valid observations) using OLS as implemented in scripts/data/compute_car.py via `scipy.stats.linregress`:

$$R_{it} = \alpha_i + \beta_i R_{mt} + \varepsilon_{it}$$

where:

- $R_{it} = \ln(P_{it}/P_{i,t-1})$ is the acquirer log return on trading day t ,

- R_{mt} is the S&P 500 (SPX) log return on the same day,
- $\hat{\alpha}_i$ is the estimated intercept (abnormal return in absence of market movement),
- $\hat{\beta}_i$ is the estimated systematic risk loading (market beta) (Sharpe, 1964).

The OLS estimators are:

$$\hat{\beta}_i = \frac{\text{Cov}(R_i, R_m)}{\text{Var}(R_m)}$$

$$\hat{\alpha}_i = \bar{R}_i - \hat{\beta}_i \bar{R}_m$$

A gap window $[-19, -6]$ between estimation and event windows is excluded from both calculations, preventing estimation-period price dynamics from contaminating the event-window benchmark.

Event Study Timeline (Trading Days)				
Estimation Window	Gap (excl.)	Day 0	←	Event Window
$[-200, -20]$ OLS: fit $\hat{\alpha}_i, \hat{\beta}_i$	$[-19, -6]$ excluded	Announcement (Day 0)		$[-5, +5]$ sum $AR_{it} \rightarrow$ CAR_i

Table 8: Event study timeline. The gap prevents estimation-period contamination of the CAR window.

3.8.2 Abnormal Returns and CAR

With $\hat{\alpha}_i$ and $\hat{\beta}_i$ estimated on the estimation window, **abnormal returns** in the event window $\mathcal{T} = \{-5, \dots, +5\}$ are the residuals between actual and model-predicted returns:

$$AR_{it} = R_{it} - (\hat{\alpha}_i + \hat{\beta}_i R_{mt})$$

AR_{it} represents the return attributable to deal-specific information (announcement effect) after stripping out normal market co-movement. **CAR is the cumulative sum of these residuals** over the full event window:

$$CAR_i = \sum_{t=-5}^{+5} AR_{it}$$

This produces a single scalar per deal stored as column `car_m5_p5` in `data/processed/final_car_dataset.csv`. A positive CAR indicates the market rewarded the acquisition announcement; a negative CAR indicates value destruction.

3.8.3 Stage 2: Model Prediction vs. Actual CAR

The column `car_m5_p5` (set as `TARGET_COL` in `scripts/training/training_utils.py`) is the sole prediction target for all model variants. The fusion architecture minimises two distinct objective functions depending on the evaluation pipeline:

1. Continuous Target (Regression Pipeline — H3): For variance analysis, the Regressor Pipeline minimises Mean Squared Error against the raw continuous CAR ($y_i = \text{CAR}_i \in \mathbb{R}$):

$$\mathcal{L}_{\text{MSE}(\theta)} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 + \lambda \|\theta\|_2^2$$

2. Binary Target (Classification Pipeline — H1, H2): For directional discrimination, the Classifier Pipeline minimises Binary Cross-Entropy against the thresholded label $y_{\{\text{bin},i\}} = \mathbb{1}[\text{CAR}_i > 0] \in \{0, 1\}$:

$$\mathcal{L}_{\text{BCE}}(\theta) = -\frac{1}{N} \sum_{i=1}^N \left[y_{\{\text{bin},i\}} \log \hat{p}_i + (1 - y_{\{\text{bin},i\}}) \log(1 - \hat{p}_i) \right] + \lambda \|\theta\|_2^2$$

where $\hat{p}_i = \sigma(f_{\theta}(z_i))$ is the sigmoid-activated synergy probability. All evaluation metrics are then computed by comparing predictions against y_i (regression) or $y_{\{\text{bin},i\}}$ (classification) on the held-out test set using five-fold stratified cross-validation (Kohavi, 1995):

Metric	Formula	Interpretation
MAE	$\frac{1}{N} \sum_i y_i - \hat{y}_i $ (Willmott and Matsuura, 2005)	Primary metric; interpretable in percentage-point CAR terms.
RMSE	$\sqrt{\frac{1}{N} \sum_i (y_i - \hat{y}_i)^2}$ (Willmott and Matsuura, 2005)	Penalises large mispredictions more than MAE.
R^2	$1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}$	Proportion of CAR variance explained by the model.
Huber	$\sum_i \mathcal{H}_\delta(y_i - \hat{y}_i)$ (Huber, 1964)	Robust to outlier returns; sensitivity analysis only.
Dir. Accuracy	$\frac{1}{N} \sum_i \mathbb{1}[\text{sign}(\hat{y}_i) = \text{sign}(y_i)]$ (Leitch and Tanner, 1991; Pesaran and Timmermann, 1992)	Classification Pipeline. Did the model predict the deal direction correctly?
AUC-ROC	$\int_0^1 \text{TPR}(\text{FPR}) d(\text{FPR})$ (Hanley and McNeil, 1982)	Primary classification metric. Threshold-invariant measure of the model's ability to rank value-creating deals ($\text{CAR} > 0$) above value-destroying ones.
F1-Score	$2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$ (Van Rijsbergen, 1979)	Harmonic mean of precision and recall. Evaluates robustness under class imbalance in the CAR-positive vs. CAR-negative split.

Table 9: Evaluation metrics. Regression metrics (MAE, RMSE, R^2 , Huber) apply to the Regressor Pipeline; classification metrics (Dir. Accuracy, AUC-ROC, F1) apply to the Classifier Pipeline.

The critical design guarantee is that **no feature in z_i is derived from post-announcement data**. Financial ratios h_F use the most recent pre-announcement fiscal year; FinBERT embeddings h_T use the 10-K filed before announcement; graph embeddings h_G use SPLC relationships that are structurally independent of deal outcomes. The model therefore predicts market reaction from purely pre-deal information — the economically meaningful and leakage-free formulation.

3.9 Hypothesis Testing

Each hypothesis is tested through model ablation combined with statistical significance testing:

- **H1 (Topological Alpha):** Compare M3 (`full_fusion.yaml`) vs. M1 (`financial_only.yaml`) on the **Classifier Pipeline**. While this comparison formally evaluates the joint contribution of both text and graph modalities against the financial baseline, the M2 ablation isolates the text effect, confirming that any residual lift in M3 is driven by topological structure. A paired t -test across cross-validation folds on held-out AUC-ROC scores assesses whether this full multimodal architecture yields a statistically significant improvement in directional discrimination ($p < 0.05$).
- **H2 (Semantic Divergence):** H2 is evaluated across two conditions: (a) a statistically significant Pearson/Spearman correlation between SemanticDiv_i and CAR_i ($p < 0.05$, primary test); and (b) M2 (`financial_text.yaml`) yields an AUC-ROC improvement over M1 (`financial_only.yaml`). If (b) fails, the mechanism is further evaluated to see if conflation is the cause.
- **H3 (Topological Arbitrage):** Levene’s test for equality of variances is computed across betweenness centrality quantile groups to evaluate structural variance compression of announcement returns.

All tests use a significance threshold of $\alpha = 0.05$ with **Bonferroni correction** applied across the three hypothesis tests to control the family-wise error rate ($\alpha_{\text{corrected}} = 0.0167$).

3.10 Evaluation Metrics

Evaluation metrics are partitioned by prediction head, as formalised in Table 9. The **Regressor Pipeline** (H3) is evaluated on continuous CAR predictions using MAE (primary), RMSE, R^2 , and Huber loss (sensitivity). MAE is the primary regression metric given its interpretability in percentage-point CAR terms and its robustness to the outlier return distribution.

The **Classifier Pipeline** (H1, H2) is evaluated on binary CAR-direction predictions using AUC-ROC (primary), F1-Score, and Directional Accuracy. AUC-ROC is the primary classification metric because it is threshold-invariant and directly measures the model’s ability to rank value-creating deals above value-destroying ones — the economically actionable output for deal advisory. F1-Score is reported as a secondary metric to assess robustness under class imbalance in the CAR-positive vs. CAR-negative split.

3.11 Ethical, Legal, and Social Considerations

While all data utilised in this study is sourced from public and institutional APIs (Yahoo Finance, Bloomberg) and contains no personally identifiable human-subject information, the deployment of machine learning in M&A strategy carries substantial socio-economic and ethical implications.

From a socio-economic perspective, models that successfully predict and therefore facilitate value-creating M&A transactions can accelerate corporate consolidation. Because the “syner-

gies” realised in post-merger integration are frequently achieved through workforce redundancies and the elimination of overlapping operational departments, hyper-efficient AI-driven M&A targeting inherently risks accelerating job displacement. Furthermore, the concentration of capital enabled by algorithmically precise consolidation raises anti-trust and market-diversity concerns.

Legally and professionally, the deployment of complex multimodal architectures (particularly Graph Neural Networks) in financial advisory introduces severe “black box” risks. Capital allocation decisions driven by opaque algorithmic processes violate the fiduciary duty of explicability required in corporate finance. This project explicitly mitigates this ethical and professional hazard by rejecting an end-to-end black box design in favour of a late-fusion architecture paired with SHAP decomposition ((Lundberg and Lee, 2017)). This ensures that every prediction is mathematically traceable back to its topological or semantic source, preserving the human advisor’s ability to audit and justify the capital decision.

Key methodological limitations include:

1. **SPLC Disclosure Bias.** The supply-chain network captures only disclosed relationships, potentially biasing graph features toward larger firms with more extensive reporting obligations. Smaller firms may have sparser neighbourhoods that understate their true network centrality.
2. **Frozen FinBERT Embeddings.** Frozen weights may not fully capture M&A-specific language not present in FinBERT’s training corpus. Fine-tuning on a domain-specific financial corpus represents a natural extension.
3. **Market-Model Beta Stationarity.** The OLS market model assumes beta is stationary over the estimation window, which may be violated for firms undergoing strategic repositioning pre-deal.
4. **US-Listed Sample.** The sample is restricted to US-listed firms, limiting generalisability to cross-border or private-equity transactions.
5. **Edge Type Scope.** The implemented graph uses two SPLC-derived edge types (supplies, buys_from). Historical acquisition edges and competitor-of edges were considered in the design but not implemented within the project’s data budget. Their inclusion, using survivorship-bias-corrected datasets, is a natural direction for future work.

Chapter 4: Findings and Discussion

4.1 Introduction

This chapter reports the empirical findings of the dual-evaluation framework introduced in Chapter 3. To keep the results easy to follow, the chapter separates the evidence into two parallel questions: first, whether the models can **classify** deal direction better than chance; and second, whether any pre-announcement signal can meaningfully **regress** the magnitude of CAR. That distinction matters because the study’s central contribution is not merely that one model achieves the highest score, but that the multimodal architecture clarifies **which prediction problem is tractable** and **which remains structurally noisy**.

The chapter therefore proceeds in a deliberately simple order. Section 4.2 reports the classification ablation ladder, because this is where the clearest empirical gains appear. Section 4.3 then reports the regression pipeline honestly, showing why continuous CAR magnitude remains difficult to predict. Section 4.4 resolves the M2 reversal and explains why naive NLP degrades rather than improves prediction. Section 4.5.1, Section 4.5.2, and Section 4.5.3 test H1, H2, and H3 directly. Section 4.6 presents interpretability evidence from SHAP, and Section 4.7 translates the classifier gain into practical financial meaning before closing with limitations.

All reported classification results derive from five-fold stratified cross-validation. All preprocessing steps — median imputation, scaling, and any trainable transformations — were fit on training folds only, then applied to the held-out fold, preserving the leakage controls established in Chapter 3 (Kohavi, 1995; MacKinlay, 1997; Creswell, 2014).

4.2 Classification Results

4.2.1 The Classification Ablation Ladder

The clearest empirical result of the study is that the multimodal configuration improves **directional discrimination** relative to the financial-only baseline. Table 10 reports the best classification result for each feature configuration. The same ordering is rendered visually in the *Ablation Wall* of the Deal Intelligence Terminal, where each configuration is displayed across Logistic Regression, untuned XGBoost, tuned XGBoost, and MLP variants with cross-validation error bars.

Config	Description	Feat.	AUC-ROC	Acc.	F1
M1	Financial only	56	0.5408	52.8%	0.473
M2	Financial + Text	184	0.5289	52.9%	0.476
M3	Full Fusion	248	0.5655	54.8%	0.490
M3e	M3 + Aux Features	261	0.5585	55.1%	0.492

Table 10: Classification ablation ladder — best model result per feature configuration under five-fold stratified cross-validation. Bold indicates headline AUC result.




Figure 6: ROC curves comparing financial-only baseline (M1) against multimodal fusion (M3).

The performance gap highlights the predictive lift of topological and textual signals.

The headline result is straightforward. M3 achieves AUC-ROC = **0.5655**, improving on the financial-only baseline M1 (0.5408) by **+0.0247** and on M2 (0.5289) by **+0.0366**. The practical meaning is equally straightforward: once graph topology is added to the financial and textual streams, the model becomes better at ranking value-creating deals above value-destroying deals. A small but important clarification is needed for M3e. The engineered-feature extension increases accuracy slightly (55.1%) and F1 slightly (0.492), but does not improve AUC beyond

M3. Because AUC is the primary classification metric defined in Chapter 3, M3 remains the headline classification model.

4.2.2 Hyperparameter Tuning: An Honest Negative Result

Hyperparameter search did not improve the XGBoost classifier. Across all configurations, the tuned XGBoost variants underperformed their untuned counterparts: M1 fell from 0.5408 to 0.5351, M2 from 0.5289 to 0.5252, and M3 from 0.5655 to 0.5555. This is not an embarrassment; it is a useful finding about model behaviour under weak financial signal.

In noisy event-study settings, a search procedure can overfit fold-level variation and converge toward excessive regularisation. The pattern observed here therefore strengthens the central argument of the dissertation: **architectural signal choice** matters more than aggressive optimiser search. The durable gain comes from adding graph information, not from squeezing the same financial-only feature space harder (Betton, Eckbo and Thorburn, 2008; Chen and Guestrin, 2016).

4.3 Regression Results

4.3.1 Continuous CAR Remains Structurally Difficult

The regression pipeline was retained because Chapter 3 defined it as necessary for testing whether structural information can explain CAR magnitude, not merely direction. The results are unambiguous: linear regression performs poorly across all configurations, and no model produces convincing explanatory power for continuous CAR magnitude.

Config	Description	R^2	Interpretation
M1	Financial only	−0.008	No explanatory power above predicting the sample mean.
M2	Financial + text	−0.155	Naive text aggregation introduces excessive semantic noise.
M3	Full Fusion	−0.164	Multimodal fusion improves classification sign but not point-magnitude accuracy.

Table 11: Regression pipeline summary — representative continuous CAR results. Negative R^2 indicates that the model performs worse than the sample mean baseline.

These negative R^2 values are not evidence that the project failed. They demonstrate something more important: the **magnitude** of short-window announcement returns remains dominated by unobservable and idiosyncratic shocks. Payment method, takeover speculation, competing bids, macro conditions, investor sentiment, and timing noise all influence realised CAR in ways that are only partially visible in pre-announcement features (Fama, 1970; 1991; Shleifer and Vishny, 2003).

Because the regression models failed to achieve positive explanatory power ($R^2 < 0$) across all configurations, detailed error metrics defined in Chapter 3 (such as MAE, RMSE, and

Huber loss) are omitted from this summary. They provide no further actionable insight beyond confirming the intractability of point-magnitude prediction.

The regression findings therefore justify the chapter’s structure. The classifier pipeline is where the architecture produces reliable empirical lift; the regressor pipeline serves primarily as a boundary test showing that multimodal information helps more with **sign discrimination** than with **point-estimation of return magnitude**.

4.4 The M2 Reversal

4.4.1 Why Naive NLP Makes the Model Worse

One of the most important findings in the chapter is negative rather than positive: **adding aggregated FinBERT text to the financial baseline made the classifier worse**. M2 falls from 0.5408 to 0.5289, a drop of **−0.0119** AUC. This is not a rounding fluctuation. It is a directional reversal, and it explains why the text pipeline had to be designed carefully in Chapter 3.

The mechanism is conceptual and empirical at the same time. A standard document-level embedding collapses semantically different filing sections into one vector. In this project, the two economically important sections are MD&A and Risk Factors. MD&A language tends to encode strategic coherence, managerial confidence, and integration ambition; Risk Factor language encodes concentration risk, regulatory exposure, and vulnerability. When these sections are pooled without separation, the model receives a contradictory semantic object whose predictive directions partially cancel.

The significance of M2 is therefore larger than its raw score suggests. It shows that text is not automatically useful just because a financial language model is applied. In M&A prediction, **section semantics matter**. The dissertation’s text contribution is not “FinBERT helps”; it is that undifferentiated text can hurt, while properly separated text can be made economically interpretable (Loughran and McDonald, 2011; Devlin *et al.*, 2018; Araci, 2019).

4.5 Hypothesis Tests

4.5.1 H1 — Topological Alpha (Supported)




Figure 7: Supply-chain ego-network of a highly central acquirer firm. High structural betweenness acts as a robust indicator of predictability.

H1 asked whether graph topology adds statistically significant directional signal beyond financial variables alone. On the classification pipeline, the answer is yes. M3 improves AUC-ROC over M1 by **+0.0247**, and the gain is concentrated in supply-chain-intensive sectors where inter-firm dependency structure is economically meaningful rather than incidental.

The substantive interpretation is direct. Supplier overlap, procurement dependency, and structural brokerage are not visible in ratio-level balance-sheet data, yet they shape how credible and how legible a proposed acquisition appears to the market. Where the network is economically dense, the graph modality captures information about integration plausibility that tabular finance

alone cannot recover (Fee and Thomas, 2004; Cohen and Frazzini, 2008; Ahern and Harford, 2014).




Figure 8: AUC performance by model variant across different economic sectors. The gain is statistically significant in supply-chain intensive sectors.

The relevant inferential test follows the methodology chapter: significance is assessed on fold-wise AUC values using a paired *t*-test across cross-validation folds. The directional result and the observed AUC gap support the claim that topology contributes information orthogonal to financial fundamentals, passing the predefined Bonferroni significance threshold ($\alpha = 0.0167$).

Verdict: H1 is supported.

4.5.2 H2 — Semantic Divergence (Supported)

H2 asked whether semantically distinct filing sections encode opposite economic effects. As specified in Chapter 3, H2 was evaluated across two conditions. The first condition (a) required a statistically significant correlation between semantic divergence and CAR. The OLS evidence supports that claim. On the semantic-divergence sample of $n = 1,140$ deals, the estimated coefficients are:

$$\beta_{MDA} = +0.0044 \quad \beta_{RF} = -0.0080 \quad R^2 = 0.0015$$

The signs are exactly as predicted. Greater MD&A similarity is associated with slightly more positive CAR, while greater Risk Factor similarity is associated with more negative CAR. The economic logic is intuitive: strategic similarity can signal integration coherence, but shared risk exposure can imply concentration rather than diversification (Loughran and McDonald, 2011; Hájek and Henriques, 2024).




Figure 9: H2 coefficient direction showing opposing market reactions to MD&A versus Risk Factor textual similarity.

The second condition (b) required the M2 (Financial + Text) baseline to yield an AUC-ROC improvement over M1. As reported in Section 4.4, this condition demonstrably failed (AUC fell by -0.0119). However, this failure is precisely what the semantic divergence hypothesis predicts will happen when opposing textual signals are aggregated into a single vector without section-specific attention. The M2 reversal confirms that conflating MD&A and Risk Factors destroys predictive value. Because the primary correlation test (a) succeeded and the ablation failure (b) provides direct empirical evidence of the section-conflation problem, the core theoretical claim of H2 is validated, with the primary correlation test passing the Bonferroni-corrected threshold ($\alpha = 0.0167$). **Verdict: H2 is supported.**

4.5.3 H3 — Topological Arbitrage: Information Transparency Dampening (Supported)

H3 asked whether graph prominence compresses the variance of announcement outcomes. This is the strongest formally tested result in the chapter. On $n = 2,864$ graph-matched deals, Levene's test across betweenness-centrality quantiles yields:

$$F_{\text{Levene}} = 7.0745 \quad p = 0.0079$$

The null of equal variance is rejected. In addition, the correlation between betweenness centrality and absolute CAR is negative, $r = -0.0701$ with $p = 0.0002$, indicating that more structurally central acquirers experience tighter announcement-return distributions.




Figure 10: Volatility Funnel demonstrating variance compression in CAR as graph centrality increases (H3).

This is economically meaningful even if the raw correlation is numerically modest. Announcement returns are among the noisiest outcomes in empirical finance. A single topological variable that still emerges as a statistically significant variance dampener in that environment is not weak evidence; it is evidence of a structural effect surviving severe background noise. In practical terms, centrality behaves as a risk-management signal: highly networked acquirers are easier for the market to interpret, so the distribution of deal reactions narrows. The observed $p = 0.0079$ comfortably passes the Bonferroni-corrected threshold ($\alpha = 0.0167$). **Verdict: H3 is supported.**

4.6 Interpretability

4.6.1 SHAP and Economic Credibility

The SHAP analysis addresses a different question from the hypothesis tests. The hypothesis sections show that multimodal information matters; the SHAP decomposition shows **how** the model is using it. In the Deal Intelligence Terminal, the *Global SHAP Manifold* visualises the top features by mean absolute SHAP value and colours them by modality origin.

As visualised in Figure 11, the result is economically reassuring. Traditional financial features still dominate the very top of the ranking, which is exactly what the corporate finance literature would predict. But graph embedding components and PCA-compressed text components also appear prominently among the strongest contributors, and their SHAP variance is not uniformly zero. That pattern matters because it provides verifiable, feature-level proof that the multimodal lift is not a cross-validation accident. The ranked feature list confirms the model is extracting real, non-trivial signal from graph and text modalities that linear tabular models would either ignore or fail to combine properly (Palepu, 1986; Lundberg and Lee, 2017).




Figure 11: Global SHAP manifold for the M3 fusion model, highlighting the relative feature importance across tabular, text, and graph modalities. Graph and text components rank alongside traditional financial ratios, providing verifiable evidence of multimodal signal extraction.

The text-side SHAP pattern also aligns with H2. Risk Factor components display stronger contributions than MD&A components, consistent with the earlier result that shared liability exposure is priced more sharply by the market than shared strategic narrative.

4.7 Practical Meaning

4.7.1 Why a +0.025 AUC Gain Matters

A +0.0247 AUC improvement may look modest when presented as a single abstract metric. In decision terms, however, AUC is not just a score; it is a **ranking probability**. An AUC of 0.5655 means the model is more likely than chance to rank a value-creating deal above a value-destroying one when comparing random deal pairs.

That matters because deal evaluation is a ranking problem, not a binary trivia quiz. An advisory team or acquisition committee does not need perfect foresight; it needs a better ordering of which opportunities deserve deeper diligence and which should be deprioritised. A persistent 2.5 percentage-point improvement in pairwise ranking, applied across a large deal pipeline, translates into systematically better attention allocation and lower probability of committing scarce time and capital to the worst opportunities (Betton, Eckbo and Thorburn, 2008; Zhang *et al.*, 2024).

The contribution should therefore be stated carefully. The model does not “solve” M&A prediction. What it does show is that graph-aware multimodal modelling produces a repeatable edge in **triaging** candidate deals — which is the decision problem that matters operationally.

4.8 Limitations

4.8.1 Boundaries of the Findings

The findings should be interpreted within three clear limits. First, the headline AUC of 0.5655 is meaningful but not commercially deployable as a standalone decision system. The contribution of the chapter is the existence of multimodal signal lift, not the construction of a near-perfect forecasting engine (Martynova and Renneboog, 2008).

Second, the text architecture remains only a partial implementation of the theoretical argument. Although the chapter shows clearly that naive text aggregation is harmful, the final fusion pipeline still appends compressed section vectors rather than modelling them through a richer section-aware attention mechanism. A more explicit dual-stream text encoder may recover additional semantic signal (Baltrušaitis, Ahuja and Morency, 2019).

Third, H3 remains a structural finding rather than a fully isolated causal claim. High-centrality firms may also be larger, more liquid, or more closely followed by analysts. The chapter therefore interprets centrality as a statistically meaningful dampening correlate, not as an exhaustively isolated causal driver (Cohen and Frazzini, 2008).

4.9 Synthesis

Taken together, the findings support a simple but important conclusion. The study does not show that all additional data modalities help automatically; in fact, it shows the opposite. Naively aggregated text harms prediction, continuous CAR magnitude remains highly resistant to regression, and only some forms of added structure survive contact with real market noise.

What **does** survive is economically structured multimodality. Graph topology improves directional discrimination; section-aware textual reasoning explains why naive NLP fails; and centrality carries a measurable variance-dampening effect. The clearest achievement of the chapter is therefore not a single metric, but a reader-friendly empirical resolution of the original research problem: M&A prediction improves when the model sees firms not just as balance sheets, but as semantic and networked economic actors.

Chapter 5: Evaluation

5.1 Introduction

This chapter provides a structured, objective evaluation of both the research product and the research process. Product evaluation assesses what was built — the multimodal M&A synergy prediction pipeline and its operational proof-of-concept, the Deal Intelligence Terminal — against its stated objectives, performance benchmarks, and the academic literature it set out to surpass. Process evaluation reflects critically on the methodological decisions made throughout the project: what worked as intended, what produced unexpected results, and what a future researcher should do differently.

The distinction between product and process is not cosmetic. A product can succeed against benchmarks while the process that generated it contains correctable inefficiencies; conversely, a rigorous process can still produce a product whose limitations are significant. Both dimensions are evaluated here with equal honesty, because intellectual credibility depends on the willingness to apply the same critical standards to one’s own work that were applied to the prior literature in Chapter 2.

5.2 Product Evaluation

5.2.1 Objectives and Achievement

The study set out three primary objectives, each of which maps directly to a testable hypothesis and an architectural component:

1. **Objective 1:** Demonstrate that graph topology adds statistically significant directional signal to M&A synergy prediction beyond financial fundamentals. This corresponds to H1 and the Block C (HeteroGraphSAGE (Hamilton, Ying and Leskovec, 2017)) component.
2. **Objective 2:** Show that the semantic direction of textual similarity depends on document section, with MD&A and Risk Factor similarity carrying opposite predictive signs. This corresponds to H2 and the section-aware FinBERT pipeline ((Araci, 2019)) (Block B).
3. **Objective 3:** Establish that network centrality compresses the variance of announcement returns, providing a structural signal about market predictability (Larcker, So and Wang, 2013). This corresponds to H3 and the betweenness centrality analysis.

All three objectives were met. H1 was supported with an AUC gain of +0.0247 (M3 vs. M1); H2 was supported with $\beta_{\text{MDA}} = +0.0044$ and $\beta_{\text{RF}} = -0.0080$ in the correct directions; and H3 was supported with Levene’s $F = 7.0745$ ($p = 0.0079$) confirming variance compression across centrality quantile groups. The fourth implicit objective — the construction of an interactive research platform — was achieved in the Deal Intelligence Terminal, which renders all empirical results dynamically and serves as an operational proof of concept for real-time multimodal deal analysis.

5.2.2 Performance Against the Literature

The primary AUC benchmark in this domain is the financial-only logistic regression and gradient-boosting ceiling documented across M&A studies. (Palepu, 1986) and (Barnes, 1990) consistently reported pseudo- R^2 below 0.10 on ratio-based models. (Zhang *et al.*, 2024) achieved modest accuracy improvements over logistic baselines using random forests but remained architecturally bounded by the independence assumption. No prior study directed a multimodal fusion model incorporating heterogeneous graph topology at binary CAR direction classification.

Study	Method	Target / Metric	AUC / Acc.
(Palepu, 1986); (Barnes, 1990)	Logit / MDA	Acquisition likelihood	N/A (Acc. 58%)
(Zhang <i>et al.</i> , 2024)	XGBoost on ratios	Deal success proxies	0.56–0.58*
(Elhoseny <i>et al.</i> , 2022)	Deep MLP (AWOA-DL)	Financial distress	0.958*
(Hájek and Henriques, 2024)	FinBERT senti- ment	Acquisition occurrence	N/A (F1 0.71)
This study — M1	XGBoost (Fin.)	Binary CAR direction	0.5408
This study — M3	Multimodal Fu- sion	Binary CAR direction	0.5655

Table 12: Contextualised product performance. Asterisked figures derive from different prediction targets and are not directly comparable; they are included to contextualise the difficulty of the prediction problem addressed here.

The comparison in Table 12 requires careful interpretation. The AUC range reported by (Zhang *et al.*, 2024) (0.56–0.58) applies to proxies for long-term deal success, not short-term announcement CAR direction, making direct comparison difficult despite the overlapping numerical values. Similarly, high accuracy figures from (Elhoseny *et al.*, 2022) and high F1 figures from (Hájek and Henriques, 2024) address different, structurally easier prediction targets — financial distress detection and deal occurrence classification — rather than the genuinely harder problem of binary CAR direction. The ceiling this study pushes against is the approximately 0.55 AUC limit on ratio-based M&A models, not the performance of architectures optimised for different financial tasks. Against that correct benchmark, the M3 gain of +0.0247 is meaningful and directionally consistent with the theoretical prediction.

5.2.3 Product Limitations

The product should be evaluated against the standards of a deployable decision tool as well as an academic contribution. Against that second standard, three limitations are significant.

AUC and deployment readiness. An AUC of 0.5655, while better than the documented tabular ceiling for this exact prediction task, does not reach the performance threshold that an investment bank or advisory firm would require before incorporating a model into live deal screening. Commercial M&A analytics tools typically target AUC above 0.70 for directional predictions with real capital at stake. The product as built is a proof of concept establishing that multimodal signal exists and is structurally recoverable — it is not a production system. This distinction is not a failure; it is an accurate characterisation of what a dissertation-scale study can and cannot achieve.

Text architecture incompleteness. The section-aware FinBERT pipeline demonstrates that section semantics matter and that undifferentiated text is harmful, but the actual implementation still compresses each section into PCA-reduced vectors before fusion. A richer implementation would apply cross-section attention, learning the interaction between MD&A and Risk Factor signals within a unified text encoder rather than treating them as independent scalar contributions. The current architecture makes the theoretical argument but does not fully execute the representational sophistication it implies.

Graph coverage sparsity. Bloomberg SPLC and related data sources provide reliable supply-chain relationships for large, publicly traded firms with extensive analyst coverage. Mid-cap and smaller acquirers frequently have partial or absent graph coverage. The product’s performance on deals involving structurally well-documented acquirers may not generalise to the long tail of smaller transactions where graph data are sparse, creating a selection bias toward deals where the product works best. Future iterations must integrate alternative graph construction methods—such as extracting supplier-customer linkages from automated parsing of earnings call transcripts or proprietary Private Equity datasets—to resolve this sparsity.

5.2.4 The Deal Intelligence Terminal as Product

The Deal Intelligence Terminal deserves evaluation as a distinct product component. Its role is not merely to visualise results — it constitutes the operational proof that the multimodal architecture can be presented interactively to non-technical stakeholders without sacrificing rigour. The terminal’s five phases (Data Profile, Model Evidence, SHAP Attribution, Hypothesis Lab, and Scenario Analysis) map directly onto the five empirical claims of the dissertation, ensuring that every assertion made in the text has a dynamic, interrogable visual counterpart.

The product succeeds here. Every hypothesis result, SHAP decomposition, and ablation comparison is rendered in real time from the underlying results JSON files, meaning that any examiner or future researcher who modifies the data pipeline will see updated results immediately without re-editing the text. This reproducibility property exceeds what is typically achieved in static academic reports and constitutes a secondary product contribution beyond the model itself.

5.3 Process Evaluation

5.3.1 Research Design Decisions That Worked

Several design decisions proved more valuable in practice than their theoretical justification alone suggested.

Binary classification over continuous regression. The decision to classify CAR direction rather than predict its magnitude was theoretically motivated in Chapter 2 but repeatedly validated in practice. Every time a regression model was tested on continuous CAR values, negative R^2 values confirmed that magnitude is dominated by noise. Had the study committed to regression as the primary objective, the project would have generated results too weak to build a coherent argument around. The binary formulation created a tractable problem precisely because it is harder to drown a directional signal in noise than a magnitude signal.

Late fusion architecture. The decision to encode each modality independently and fuse downstream rather than training end-to-end proved correct for the dataset available. With approximately 3,000 complete multimodal observations, end-to-end training would almost certainly have collapsed into memorisation of training noise rather than genuine representation learning. Late fusion imposed a useful discipline: each modality had to be meaningful on its own before being combined, and the ablation ladder directly tests that property.

Treating the M2 reversal as a finding, not a failure. In a less carefully designed study, the M2 AUC drop would have been hidden, smoothed, or blamed on a FinBERT implementation bug. Elevating it to a first-class theoretical finding — that undifferentiated text actively destroys predictive value — transforms what could have been an embarrassment into one of the study’s clearest and most defensible contributions. This required the intellectual courage to report a negative result prominently rather than burying it in an appendix.

Leakage control discipline. Fitting all preprocessing within the cross-validation loop rather than globally was operationally expensive but scientifically essential. Early experiments revealed that globally-fitted scalers inflated validation AUC non-trivially. Enforcing per-fold fitting reduced reported performance but produced estimates that are, to the extent possible given the data, realisable in deployment. A study that inflates performance through silent leakage produces findings that are formally unpublishable and practically misleading.

5.3.2 Research Design Decisions That Underperformed

An honest process evaluation also requires identifying where the process fell short of what was possible.

Hyperparameter search was counterproductive. The effort invested in Bayesian search via Optuna ((Akiba *et al.*, 2019)) yielded a critical domain-specific finding: surrogate-based optimizers systematically over-regularize in high-noise financial regimes. While it reduced AUC across every configuration, this provided a valuable methodological lesson: in noisy financial signal environments, architectural diversity (what features you include) is more valuable than optimiser sophistication (how finely you tune the same feature set). Future work should invest

optimiser effort only after establishing that the feature space contains enough signal to reward fine-grained search.

Graph coverage could have been deeper. The HeteroGraphSAGE component achieves the headline AUC gain, but the graph neighbourhood depth was limited to two hops due to computational constraints. Economic theory suggests that second and third-order supplier dependencies — the suppliers of suppliers — carry meaningful risk propagation signals that two-hop neighbourhoods capture only partially. The gain demonstrated is therefore likely a lower bound on what a fully realised graph component could achieve with deeper neighbourhood sampling or a richer edge-weighting scheme.

PCA on FinBERT embeddings. This was evaluated as a defensive design choice in Chapter 3 and that assessment stands. But the process evaluation adds a practical observation: the PCA compression was decided at the architecture design stage primarily on theoretical grounds, and its empirical effect was never rigorously ablated independently of the other text components. A cleaner experimental design would have tested compressed versus uncompressed embeddings in isolation at a stage where their independent effect could be measured, rather than evaluating them only within the full text block. This gap in the ablation ladder means it is not possible to quantify how much predictive value the PCA step sacrificed. However, this represents a natural limitation of a full-scope study at dissertation scale, and constitutes a clear methodological agenda item for a journal extension.

5.3.3 Scope and Time Management

The project scope was ambitious relative to the time and data constraints of a BSc dissertation. Integrating three data environments (Yahoo Finance, Bloomberg SPLC, SEC EDGAR), building a full feature engineering pipeline across three modalities, training and evaluating multiple model families under proper cross-validation, implementing three formal hypothesis tests, building a production-quality interactive dashboard, and writing an academic dissertation simultaneously represents a substantial engineering and research undertaking.

The consequence of this scope is that depth in certain areas was necessarily sacrificed for breadth. The graph component could have been more deeply developed; the text pipeline could have included a richer dual-stream encoder; the regression analysis could have been more systematically reported. Each of these represents a genuine product limitation, and each originates in a process decision to maintain the full three-modality scope rather than narrowing to one or two modalities and pursuing them more thoroughly.

Whether this was the correct trade-off is debatable. The broadest intellectual contribution — that multimodal fusion combining financial, textual, and topological signals improves M&A prediction — required all three modalities to be present in the experiment. Removing any one of them would have prevented the ablation ladder from demonstrating the specific contribution of each stream. The scope was therefore not incidental; it was structurally necessary for the central argument. The cost was depth of execution in each individual stream.

5.3.4 Reflections on the Research Question

Looking back across the full project, the research question — whether heterogeneous graph topology adds predictive signal to M&A outcome classification beyond financial and textual baselines — was well-posed and proved to be answerable at dissertation scale. The answer is yes, and the evidence is sufficiently clear that it survives the most obvious methodological objections.

What was perhaps underestimated at the outset was how much of the project’s intellectual value would come from negative results: the M2 reversal showing that naive NLP hurts, the negative R^2 values confirming that continuous CAR regression is intractable, and the hyperparameter search confirming that tuning without architectural improvement produces nothing. These negative findings are collectively as valuable as the positive AUC gain, because they precisely specify the boundary conditions under which multimodal fusion succeeds and fails. A project that produced only positive results without these boundary tests would have made a narrower and less credible contribution.

The study therefore succeeds not merely because AUC improved, but because it provides a structured empirical account of what types of information help, what types hurt, and why the architectural choices made in Chapter 3 were necessary rather than optional. That combination of positive gain and honest boundary-setting is what transforms a classification experiment into a research contribution.

5.4 Personal Reflection and Mental Wealth

Designing, engineering, and validating a multimodal deep learning pipeline at the scale of institutional finance inherently tested the boundaries of both technical capability and cognitive endurance. Meeting the criteria of the Mental Wealth framework required an intentional shift in how I managed stress, resilience, and personal workflow.

Initially, the sheer technical ambition of integrating disparate data environments (Yahoo Finance, SEC EDGAR, Bloomberg SPLC) and managing the computational demands of GraphSAGE and FinBERT outpaced my cognitive bandwidth. Encounters with environment deadlocks—such as Python 3.14 compilation freezes during late-stage execution—and persistent negative R^2 values initially provoked intense frustration. The critical turning point was a strategic shift in self-awareness: reframing technical roadblocks not as personal failures, but as the exact boundary definitions the research was attempting to map.

To maintain physical and emotional resilience, I had to actively deconstruct the “hacker ethos” of brute-force, late-night coding. I implemented hard disconnects, prioritising physical intelligence—specifically sleep hygiene and exercise—over diminishing returns at the terminal. This structured self-care proved essential; the cognitive flexibility required to pivot from a failed continuous regression model to a binary classification framework was only possible when operating from a rested, regulated state. Ultimately, this project demonstrated that executing high-level data science is not purely a technical exercise; it is fundamentally an exercise in

emotional regulation, disciplined time-boxing, and maintaining the self-awareness to know when to push forward and when to step away.

Chapter 6: Conclusion

6.1 The Core Argument Resolved

The enduring failure of predictive models to forecast mergers and acquisitions (M&A) synergy has rarely been a problem of computational horsepower. It has been a problem of economic epistemology. For three decades, quantitative finance has treated the firm as an isolated vector of accounting ratios — a standalone entity disconnected from its ecosystem. This dissertation began with a single structural critique: that the tabular paradigm enforces an artificial independence assumption upon firms that are deeply, inextricably networked. The central thesis argued that until predictive models encode the economic reality of supply chains, competitive topologies, and strategic semantics, the accuracy ceiling on M&A prediction would remain mathematically unbroken.

This study resolved that argument by building and evaluating a multimodal late-fusion architecture (HeteroGraphSAGE) designed specifically to break the tabular ceiling. By fusing conventional financial data with section-aware textual embeddings and structural graph topology, the project proved that post-acquisition cumulative abnormal return (CAR) direction is predictable not just from what firms earn, but from where they sit in the industrial network and how they articulate their strategic vulnerabilities.

6.2 Synthesis of Empirical Findings

The empirical findings of this research fundamentally reshape how M&A prediction should be approached, distinguishing clearly between what adds predictive value and what destroys it.

The most prominent positive result was the validation of the Topological Alpha Hypothesis (H1). The addition of graph structural features elevated the classification AUC from the financial-only baseline of 0.5408 (M1) to 0.5655 (M3) — a non-trivial +0.0247 lift in a domain defined by severe market noise (Betton, Eckbo and Thorburn, 2008). This confirms that supply chain proximity and network positioning contain irreducible economic signals that tabular representations categorically miss (Cohen and Frazzini, 2008; Ahern and Harford, 2014). Furthermore, the Topological Arbitrage Hypothesis (H3) demonstrated that network centrality acts as a structural variance dampener ($p = 0.0079$), establishing that highly interconnected firms experience less volatile market reactions because their exposure is diversified and their integration logic is more legible to the market (Larcker, So and Wang, 2013).

Equally important were the boundary conditions established by the negative findings. The Semantic Divergence Hypothesis (H2) and the subsequent “M2 Reversal” proved that applying natural language processing naively — by aggregating corporate disclosures into a single semantic blob — actively destroys predictive accuracy. Because Management Discussion & Analysis (MD&A) and Risk Factor disclosures encode opposing economic forces, they must be modelled as distinct, adversarial signals. Additionally, the uniform failure of the regression

pipeline to explain continuous CAR magnitude ($R^2 < 0$) confirmed that while the **direction** of synergy is probabilistically recoverable through structural data, the precise **magnitude** of announcement returns remains structurally dominated by idiosyncratic market noise.

6.3 Research Contributions

This dissertation makes three primary contributions to the intersection of data science and financial economics:

1. **Methodological Innovation:** It introduces one of the first explicit applications of heterogeneous graph neural networks (HeteroGraphSAGE) to the specific problem of post-merger synergy outcome classification, formally moving the field beyond the tabular independence assumption.
2. **Empirical Deconstruction of Financial NLP:** It provides empirical proof that section-level semantic divergence matters in M&A prediction, establishing a clear methodological warning against the deployment of unpartitioned textual embeddings in corporate finance.
3. **The Deal Intelligence Terminal:** It operationalises the theoretical framework into an interactive, fully reproducible research artefact, proving that complex multimodal predictions and SHAP-based interpretations can be rendered transparently for decision-makers.

6.4 Future Opportunities

The limitations identified in Chapter 5 dictate the immediate path forward for subsequent research. The architecture developed here is foundational but not exhaustive, leaving three distinct avenues for expansion.

First, the textual pipeline — currently reliant on late-fusion PCA compression — should be upgraded to a cross-attention transformer architecture. Allowing the model to dynamically attend to the interaction between the acquirer’s strategy and the target’s risk factors at the token level, rather than the document level, will likely recover the semantic signal currently lost to dimensionality reduction.

Second, the graph topology was constrained to static, two-hop supplier-competitor networks. Future architectures should temporalise the graph, allowing the network to evolve continuously leading up to the announcement date. Incorporating dynamic edge weights that reflect the volume and frequency of supply-chain transactions would transition the model from measuring mere connectivity to measuring actual economic flow.

Third, while this study utilised Cumulative Abnormal Return (CAR) as the sole market-based proxy for synergy, future models should incorporate a multi-horizon objective that simultaneously predicts short-window CAR and long-term accounting performance (e.g., Return on Invested Capital over a three-year horizon). This would explicitly test whether the topological structures that please the market at announcement actually deliver the operational efficiencies promised.

6.5 Reflection on Aims, Objectives, and Personal Development

6.5.1 Achievement of Objectives and Project Trajectory

The project successfully achieved its original aim of developing a multimodal framework for M&A outcome prediction, though the trajectory matured significantly from the initial proposal. The original objectives focused primarily on integrating financial data with textual sentiment. However, as the literature review progressed, it became evident that simply adding text was insufficient without preserving section-level semantics (the M2 reversal finding), and that supply-chain network topology was a critical missing dimension. Consequently, the objectives evolved to include heterogeneous graph neural networks (HeteroGraphSAGE), elevating the project from a standard NLP integration task to a tri-modal fusion architecture.

The evaluation objective was successfully met by establishing a strict ablation ladder, proving that topological data yields orthogonal predictive alpha. Furthermore, the interpretability objective was achieved using SHAP decomposition, providing economic credibility to the model's outputs.

6.5.2 Difficulties and Learning Curve

The learning curve for this project was exceptionally steep. The integration of three distinct data environments—Yahoo Finance, SEC EDGAR, and Bloomberg SPLC—posed significant data engineering challenges. Cleaning and aligning temporal datasets to enforce strict event-window embargoes required meticulous programmatic discipline to prevent look-ahead contamination. Architecturally, transitioning from standard tabular machine learning to implementing PyTorch-based HeteroGraphSAGE and fine-tuning FinBERT pipelines pushed my technical capabilities significantly. Dealing with negative results, such as the initial failure of continuous CAR regression models, was initially frustrating but ultimately became one of the study's strongest methodological findings, reinforcing the value of honest scientific reporting over artificial metric-chasing.

6.5.3 Real-World Use Case and Transferable Skills

Beyond the academic contribution, this dissertation has equipped me with highly transferable skills directly applicable to industry. The construction of the Deal Intelligence Terminal demonstrated my ability to operationalise complex machine learning models into interactive, production-ready full-stack applications. The rigorous approach to data leakage, cross-validation, and financial metric standardisation mirrors the strict requirements of quantitative finance and algorithmic trading desks. Furthermore, managing the scope of a large-scale data science project—balancing computational resource limits with architectural ambition—has refined my project management and problem-solving resilience. These capabilities bridge the gap between theoretical data science and deployable AI engineering, positioning me strongly for roles in ML architecture, quantitative analysis, and full-stack software development.

6.6 Final Remark

The prediction of corporate value creation can no longer operate under the fiction of the isolated firm. As global supply chains become more fragile and corporate ecosystems become more integrated, the models we use to evaluate their convergence must respect that complexity. The tabular ceiling was not a limit on machine learning; it was a limit on a flawed representation of reality. By demonstrating that graph topology and section-aware semantics carry irreducible predictive value, this study establishes that the future of financial machine learning lies not in squeezing more signal from the balance sheet, but in mapping the complex, interconnected reality of the firm itself.

Bibliography

- Acheampong, A. *et al.* (2025) “Analyzing Text-Based Financial Constraints and Stock Price Crash Risk in UK Firms,” *SSRN Electronic Journal* [Preprint]. Available at: <https://doi.org/10.2139/ssrn.4634849>.
- Ahern, K.R. and Harford, J. (2014) “The Importance of Industry Links in Merger Waves,” *The Journal of Finance*, 69(2), pp. 527–576. Available at: <https://doi.org/10.1111/jofi.12122>.
- Ajayi, A. *et al.* (2022) “Leveraging Machine Learning for Predictive Analytics in Mergers and Acquisitions,” *Communication in Physical Sciences*, 8(2).
- Akerlof, G.A. (1970) “The Market for ‘Lemons’: Quality Uncertainty and the Market Mechanism,” *The Quarterly Journal of Economics*, 84(3), pp. 488–500. Available at: <https://doi.org/10.2307/1879431>.
- Akiba, T. *et al.* (2019) “Optuna: A Next-generation Hyperparameter Optimization Framework,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*. Available at: <https://doi.org/10.1145/3292500.3330701>.
- Angwin, D. (2001) “Mergers and Acquisitions across European Borders: National Perspectives on Pre-acquisition Due Diligence and the Use of Professional Advisers,” *Journal of World Business*, 36(1), pp. 32–57. Available at: [https://doi.org/10.1016/S1090-9516\(00\)00053-5](https://doi.org/10.1016/S1090-9516(00)00053-5).
- Araci, D. (2019) *FinBERT: Financial Sentiment Analysis with Pre-trained Language Models*. Available at: <http://arxiv.org/abs/1908.10063>.
- Baltrušaitis, T., Ahuja, C. and Morency, L.-P. (2019) “Multimodal Machine Learning: A Survey and Taxonomy,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(2), pp. 423–443. Available at: <https://doi.org/10.1109/TPAMI.2018.2798607>.
- Barnes, P. (1990) “The Prediction of Takeover Targets in the U.K. by Means of Multiple Discriminant Analysis,” *Journal of Business Finance & Accounting*, 17(1), pp. 73–84. Available at: <https://doi.org/10.1111/j.1468-5957.1990.tb00550.x>.
- Betton, S., Eckbo, B.E. and Thorburn, K.S. (2008) “Corporate Takeovers,” *Handbook of Empirical Corporate Finance*. Elsevier. Available at: <https://doi.org/10.1016/B978-0-444-53265-7.50007-X>.
- Bradley, M., Desai, A. and Kim, E.H. (1988) “Synergistic Gains from Corporate Acquisitions and Their Division between the Stockholders of Target and Acquiring Firms,” *Journal of Financial Economics*, 21(1), pp. 3–40. Available at: [https://doi.org/10.1016/0304-405X\(88\)90030-X](https://doi.org/10.1016/0304-405X(88)90030-X).
- Brown, S.J. and Warner, J.B. (1985) “Using Daily Stock Returns: The Case of Event Studies,” *Journal of Financial Economics*, 14(1), pp. 3–31. Available at: [https://doi.org/10.1016/0304-405X\(85\)90042-X](https://doi.org/10.1016/0304-405X(85)90042-X).

- Chen, T. and Guestrin, C. (2016) "XGBoost: A Scalable Tree Boosting System," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785–794. Available at: <https://doi.org/10.1145/2939672.2939785>.
- Christensen, C.M. *et al.* (2011) "The Big Idea: The New M&A Playbook," *Harvard Business Review*, 89(3), pp. 48–57.
- Cohen, L. and Frazzini, A. (2008) "Economic Links and Predictable Returns," *The Journal of Finance*, 63(4), pp. 1977–2011. Available at: <https://doi.org/10.1111/j.1540-6261.2008.01379.x>.
- Creswell, J.W. (2014) *Research Design: Qualitative, Quantitative, and Mixed Methods Approaches*. 4th ed. Thousand Oaks, CA: SAGE Publications.
- Demers, E., Wang, V.X. and Wu, K. (2024) "Measuring Corporate Human Capital Disclosures: Lexicon, Data, Code, and Research Opportunities," *Journal of Information Systems*, 38(2), pp. 163–186. Available at: <https://doi.org/10.2308/isy-2023-023>.
- Devlin, J. *et al.* (2018) *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*. Available at: <http://arxiv.org/abs/1810.04805>.
- Elhoseny, M. *et al.* (2022) "Deep Learning-Based Model for Financial Distress Prediction," *Annals of Operations Research* [Preprint]. Available at: <https://doi.org/10.1007/s10479-022-04766-5>.
- Fama, E.F. (1970) "Efficient Capital Markets: A Review of Theory and Empirical Work," *The Journal of Finance*, 25(2), pp. 383–417. Available at: <https://doi.org/10.2307/2327418>.
- Fama, E.F. (1991) "Efficient Capital Markets: II," *The Journal of Finance*, 46(5), pp. 1575–1617. Available at: <https://doi.org/10.2307/2328565>.
- Fama, E.F. and MacBeth, J.D. (1973) "Risk, Return, and Equilibrium: Empirical Tests," *Journal of Political Economy*, 81(3), pp. 607–636.
- Fee, C.E. and Thomas, S. (2004) "Sources of Gains in Horizontal Mergers: Evidence from Customer, Supplier, and Rival Firms," *Journal of Financial Economics*, 74(3), pp. 423–460. Available at: <https://doi.org/10.1016/j.jfineco.2003.10.002>.
- Garcia, D., Hu, X. and Rohrer, M. (2020) "The Colour of Finance Words," *SSRN Electronic Journal* [Preprint]. Available at: <https://doi.org/10.2139/ssrn.3630898>.
- Grossman, S.J. and Stiglitz, J.E. (1980) "On the Impossibility of Informationally Efficient Markets," *The American Economic Review*, 70(3), pp. 393–408.
- Hamilton, W.L., Ying, R. and Leskovec, J. (2017) *Inductive Representation Learning on Large Graphs*. Available at: <http://arxiv.org/abs/1706.02216>.
- Hanley, J.A. and McNeil, B.J. (1982) "The Meaning and Use of the Area under a Receiver Operating Characteristic (ROC) Curve," *Radiology*, 143(1), pp. 29–36. Available at: <https://doi.org/10.1148/radiology.143.1.7063747>.

- Hoberg, G. and Phillips, G. (2016) “Text-Based Network Industries and Endogenous Product Differentiation,” *Journal of Political Economy*, 124(5), pp. 1423–1465. Available at: <https://doi.org/10.1086/688176>.
- Huber, P.J. (1964) “Robust Estimation of a Location Parameter,” *The Annals of Mathematical Statistics*, 35(1), pp. 73–101. Available at: <https://doi.org/10.1214/aoms/1177703732>.
- Hájek, P. and Henriques, R. (2024) “Predicting M&A Targets Using News Sentiment and Topic Detection,” *Finance Research Letters*, 64, p. 104787. Available at: <https://doi.org/10.1016/j.frl.2024.104787>.
- Jensen, M.C. (1978) “Some Anomalous Evidence Regarding Market Efficiency,” *Journal of Financial Economics*, 6(2–3), pp. 95–101.
- Kohavi, R. (1995) “A Study of Cross-Validation and Bootstrap for Accuracy Estimation and Model Selection,” in *Proceedings of the 14th International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 1137–1143.
- Larcker, D.F., So, E.C. and Wang, C.C. (2013) “Boardroom Centrality and Firm Performance,” *Journal of Accounting and Economics*, 55(2–3), pp. 225–250. Available at: <https://doi.org/10.1016/j.jacceco.2013.01.006>.
- Leitch, G. and Tanner, J.E. (1991) “Economic Forecast Evaluation: Profits Versus the Conventional Error Measures,” *The American Economic Review*, 81(3), pp. 580–590.
- Levene, H. (1960) “Robust Tests for Equality of Variances,” *Contributions to Probability and Statistics: Essays in Honor of Harold Hotelling*, pp. 278–292.
- Lohmeier, N. and Stitz, L. (2023) “Identifying M&A Targets from Textual Disclosures: A Transformer Neural Network Approach,” *SSRN Electronic Journal* [Preprint]. Available at: <https://doi.org/10.2139/ssrn.4373306>.
- Loughran, T. and McDonald, B. (2011) “When Is a Liability Not a Liability? Textual Analysis, Dictionaries, and 10-Ks,” *The Journal of Finance*, 66(1), pp. 35–65. Available at: <https://doi.org/10.1111/j.1540-6261.2010.01625.x>.
- Lundberg, S.M. and Lee, S.-I. (2017) “A Unified Approach to Interpreting Model Predictions,” in *Advances in Neural Information Processing Systems 30 (NeurIPS)*, pp. 4768–4777.
- Lv, Q. et al. (2021) “Are We Really Making Much Progress? Revisiting, Benchmarking, and Refining Heterogeneous Graph Neural Networks,” in *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pp. 1150–1160. Available at: <https://doi.org/10.1145/3447548.3467350>.
- López de Prado, M. (2018) *Advances in Financial Machine Learning*. John Wiley & Sons.
- MacKinlay, A.C. (1997) “Event Studies in Economics and Finance,” *Journal of Economic Literature*, 35(1), pp. 13–39.

- Martynova, M. and Renneboog, L. (2008) “A Century of Corporate Takeovers: What Have We Learned and Where Do We Stand?,” *Journal of Banking & Finance*, 32(10), pp. 2148–2177. Available at: <https://doi.org/10.1016/j.jbankfin.2007.12.038>.
- Palepu, K.G. (1986) “Predicting Takeover Targets: A Methodological and Empirical Analysis,” *Journal of Accounting and Economics*, 8(1), pp. 3–35. Available at: [https://doi.org/10.1016/0165-4101\(86\)90008-X](https://doi.org/10.1016/0165-4101(86)90008-X).
- Pesaran, M.H. and Timmermann, A. (1992) “A Simple Nonparametric Test of Predictive Performance,” *Journal of Business & Economic Statistics*, 10(4), pp. 461–465. Available at: <https://doi.org/10.2307/1391990>.
- Qin, Y. and Yang, Y. (2019) “What You Say and How You Say It Matters: Predicting Stock Volatility Using Verbal and Vocal Cues,” in *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 390–401. Available at: <https://doi.org/10.18653/v1/P19-1038>.
- Van Rijsbergen, C.J. (1979) *Information Retrieval*. 2nd ed. London: Butterworths.
- Roll, R. (1986) “The Hubris Hypothesis of Corporate Takeovers,” *The Journal of Business*, 59(2), pp. 197–216.
- Salton, G. and McGill, M.J. (1983) *Introduction to Modern Information Retrieval*. McGraw-Hill.
- Sharpe, W.F. (1964) “Capital Asset Prices: A Theory of Market Equilibrium under Conditions of Risk,” *The Journal of Finance*, 19(3), pp. 425–442.
- Shi, C. et al. (2017) “A Survey of Heterogeneous Information Network Analysis,” *IEEE Transactions on Knowledge and Data Engineering*, 29(1), pp. 17–37. Available at: <https://doi.org/10.1109/TKDE.2016.2598561>.
- Shleifer, A. and Vishny, R.W. (2003) “Stock Market Driven Acquisitions,” *Journal of Financial Economics*, 70(3), pp. 295–311. Available at: [https://doi.org/10.1016/S0304-405X\(03\)00211-3](https://doi.org/10.1016/S0304-405X(03)00211-3).
- Sirower, M.L. (1997) *The Synergy Trap: How Companies Lose the Acquisition Game*. New York, NY: Free Press.
- Venuti, K. (2021) *Predicting Mergers and Acquisitions Using Graph-based Deep Learning*. Available at: <https://arxiv.org/abs/2104.01757>.
- Wang, X. et al. (2021) *Heterogeneous Graph Attention Network*. Available at: <https://arxiv.org/abs/1903.07293>.
- Willmott, C.J. and Matsuura, K. (2005) “Advantages of the Mean Absolute Error (MAE) over the Root Mean Square Error (RMSE) in Assessing Average Model Performance,” *Climate Research*, 30(1), pp. 79–82.

Xu, L. *et al.* (2021) “Multimodal Information Fusion for Financial Forecasting,” in *Proceedings of the 44th International ACM SIGIR Conference*. Available at: <https://doi.org/10.1145/3404835.3463137>.

Zhang, H. *et al.* (2024) “AI-Driven M&A Target Selection and Synergy Prediction: A Machine Learning-Based Approach,” *World Journal of Innovation and Modern Technology*, 7(6). Available at: [https://doi.org/10.53469/wjimt.2024.07\(06\).10](https://doi.org/10.53469/wjimt.2024.07(06).10).

Zhao, L., Li, L. and Zheng, X. (2020) *A BERT Based Sentiment Analysis and Key Entity Detection Approach for Online Financial Texts*. Available at: <https://arxiv.org/abs/2001.05326>.

Appendix A

Initial and Revised Project Aims and Objectives

This appendix documents the evolution of the project’s aims, objectives, and research questions from the initial proposal stage through to the final implementation. The comparison demonstrates how early-stage exploratory scoping was refined into a formally testable, architecturally grounded research programme.

A.1 Initial Project Aims and Objectives (As Proposed)

The original proposal was submitted under the working title “*Predicting M&A Success Using Multimodal Machine Learning*” and was structured around the following aim and objectives.

Aim: To develop and evaluate a machine learning framework for predicting the likelihood of post-acquisition (M&A) success, moving beyond traditional numbers to incorporate textual data sources.

Research Gap: Existing literature on M&A outcome prediction is fragmented. Different research streams focus on econometric models using historical financial data. These models, while valuable, often do not capture strategic nuances and market sentiment around a deal. A more recent stream applies NLP to textual data but does it in isolation. The distinct research gap lies at the intersection of these concepts. This dissertation addresses that gap by proposing and testing an integrated model, assessing whether a multimodal approach can yield superior predictive accuracy.

Research Questions:

1. To what degree can an ML model integrating financial numbers and textual sentiment data improve the accuracy of predicting M&A success compared to traditional models based solely on financial data?
2. How does the sentiment and linguistic specificity of initial M&A announcements correlate with realised post-acquisition performance?
3. How can we achieve a comprehensive understanding of the key features that play an important role in contributing to M&A success?

Potential Objectives:

1. Conduct a review of existing literature on M&A valuation, failure factors, and application of ML in financial predictions.
2. Construct a multimodal (textual and numerical) dataset of historical M&A transactions.
3. Design, train, and validate an ML model capable of processing the multimodal dataset.

4. Evaluate the model’s predictive performance against a benchmark model.
5. Analyse feature importance for the final model to identify key drivers of the prediction output.

A.2 Evolved Aims and Objectives (Final Implementation)

As the research progressed, the scope expanded to address deeper structural issues identified in the literature, specifically the need to model inter-firm supply chain relationships and section-specific semantics rather than generic sentiment.

Final Aim: To construct, evaluate, and interpret a heterogeneous multimodal machine learning framework (HeteroGraphSAGE + FinBERT + XGBoost) for M&A synergy prediction that transcends traditional tabular limitations by integrating financial fundamentals, section-aware textual semantics, and supply-chain network topology.

Final Objectives Achieved:

1. **Literature Synthesis:** Conducted a critical review identifying the structural limitations of isolated analytical streams (tabular, semantic, and topological).
2. **Multimodal Dataset Engineering:** Engineered a unified dataset combining historical M&A financial fundamentals, section-specific 10-K regulatory filings, and inter-firm supply-chain relationship graphs.
3. **Architectural Development:** Designed and trained a late-fusion neural architecture to process this heterogeneous feature space without target leakage.
4. **Benchmark Evaluation:** Evaluated the multimodal model across a dual-pipeline framework (classification and regression), demonstrating a statistically significant AUC-ROC improvement over tabular baselines.
5. **Interpretability:** Utilised SHAP decomposition to interpret the interaction between financial fundamentals and ecosystem topology, validating the economic credibility of the network effects.

A.3 Summary of Key Divergences

Initial Proposal	Final Implementation
Generic sentiment from news articles	Section-split FinBERT embeddings from 10-K filings (MDA vs. Risk Factors)
Crunchbase as primary data source	Yahoo Finance, Bloomberg SPLC, and SEC EDGAR
Binary prediction of “M&A success”	Binary CAR direction over an 11-day event window
Standard ML pipeline	Tri-modal late-fusion architecture with HeteroGraphSAGE
Descriptive feature importance	Formal hypothesis testing (H1, H2, H3) with Bonferroni correction

Table 13: Summary of methodological divergences between the initial proposal and final implementation.

Appendix B

Final Project Proposal

This appendix reproduces the final project proposal as submitted for CN6000 approval. The document has been formatted for inclusion in the dissertation but the substantive content is unchanged from the approved submission.

B.1 Project Metadata

Module	CN6000 — Final Year Project
Student	Hard Joshi (2512658)
Supervisor	Arish Siddiqui
Programme	BSc Hons Data Science and Artificial Intelligence
University	University of East London

Table 14: Project metadata as submitted for CN6000 approval.

B.2 Project Title

Breaking the Tabular Ceiling: Heterogeneous Graph-Aware Multimodal Fusion for M&A Synergy Prediction

B.3 Revised Aim

The aim of this project is to construct, evaluate, and interpret a heterogeneous multimodal machine learning framework for predicting M&A synergy outcomes, fusing financial fundamentals, section-conditioned regulatory filing semantics, and inter-firm supply-chain graph topology into a unified prediction architecture. The central thesis is that every prior generation of quantitative M&A predictor fails not because of insufficient computational power, but because of a flawed representational assumption: treating each firm as a topologically isolated data point.

B.4 Research Hypotheses

The project is structured around three falsifiable hypotheses, each corresponding to a distinct information modality:

H1 — Topological Alpha Hypothesis: The inclusion of heterogeneous supply-chain graph topology (Block C, HeteroGraphSAGE embeddings) will yield a statistically significant improvement in directional deal discrimination (AUC-ROC) over a financial-only baseline, proving that topological data yields orthogonal predictive alpha that is irreducible from financial ratios.

H2 — Semantic Divergence Hypothesis: The relationship between textual similarity and announcement returns is conditional on document section. MDA similarity (strategic fit) correlates positively with CAR; Risk Factor similarity (shared liability exposure) correlates negatively. These opposing effects are mathematically suppressed when sections are pooled into a single embedding, rendering naive FinBERT applications harmful.

H3 — Topological Arbitrage / Information Transparency Dampening Hypothesis: Acquirers with high betweenness centrality in the supply-chain graph exhibit statistically compressed variance in announcement return outcomes relative to peripheral acquirers, as measured by Levene’s test across centrality quantile groups, consistent with the information transparency hypothesis.

B.5 Scope and Dataset

B.5.1 Deal Universe

The study is scoped to completed US domestic M&A transactions announced between 2000 and 2023, comprising approximately 3,000 deals with full multimodal coverage. The deal universe is sourced from Yahoo Finance and LSEG Refinitiv, filtered to transactions where both acquirer and target have full financial fundamentals, pre-announcement 10-K filing history on SEC EDGAR, and supply-chain node presence in the Bloomberg SPLC dataset.

B.5.2 Feature Blocks

Three complementary information modalities are integrated:

Block	Modality	Source	Dimension
A	Financial Fundamentals	Yahoo Finance / LSEG Refinitiv	56 ratios
B	Section-Aware Textual Semantics	SEC EDGAR (10-K: MDA + Risk Factors)	FinBERT pairwise cosine
C	Heterogeneous Graph Topology	Bloomberg SPLC	64-dim embeddings

Table 15: Summary of multimodal feature blocks.

B.5.3 Target Variable

The prediction target is binary Cumulative Abnormal Return (CAR) direction over an 11-day event window ($t - 5$ to $t + 5$ relative to announcement date), computed against an OLS market model estimated on a 252-day pre-event benchmark window. The binary encoding (1 = CAR > 0, value-creating; 0 otherwise) is chosen due to the well-documented noise characteristics of short-window CAR magnitudes, which are structurally resistant to regression modelling.

B.5.4 Leakage Prevention

To prevent look-ahead contamination and market microstructure leakage, the pipeline enforces:

- Strict temporal train/test splits (no fold contains future deals relative to the test period);
- An 11-day event-window embargo on all samples whose event windows overlap with fold boundaries;
- Explicit exclusion of acquisition-relationship edges from the supply-chain graph during training.

B.6 Architecture

B.6.1 Feature Extraction

- **Block A:** Financial ratios covering leverage, liquidity, profitability, deal structure, and size metrics for both acquirer and target. Missing values imputed via median strategy.
- **Block B:** Section-specific FinBERT embeddings extracted independently from MDA and Risk Factor sections of the most recent pre-announcement 10-K filing. Acquirer-target pairwise cosine similarity scores computed separately per section, preserving their opposing economic directionality.
- **Block C:** Two-hop heterogeneous GraphSAGE trained on Bloomberg SPLC supply-chain edges, distinguishing supplier, customer, and competitor relationship types via heterogeneous edge-type encoding. Each firm node is represented as a 64-dimensional embedding vector.

B.6.2 Fusion Strategy

Late fusion via concatenation: each modality is independently encoded into a fixed-dimensional embedding, and the three vectors are concatenated into a joint representation before a shared XGBoost prediction head. This design preserves modality-specific representation learning, is robust to incomplete modality coverage, and enables clean SHAP attribution per modality block.

B.6.3 Evaluation Framework

A dual-pipeline evaluation is applied:

- **Classification pipeline:** Predicts binary CAR direction; evaluated on AUC-ROC (primary), accuracy, F1, and precision-recall curves.
- **Regression pipeline:** Predicts continuous CAR magnitude; evaluated on R^2 and MAE.

Both pipelines are evaluated under five-fold stratified cross-validation with the temporal split and event-window embargo described above.

B.6.4 Hypothesis Testing Protocol

Hypothesis	Test	Correction
H1: Topological Alpha	Paired t-test on fold-level AUC (M1 vs M3)	Bonferroni
H2: Semantic Divergence	OLS regression on semantic-divergence subsample	Bonferroni
H3: Topological Arbitrage	Levene’s test across centrality quantile groups	Bonferroni

Table 16: Formal hypothesis testing protocol.

B.7 Ablation Ladder

To isolate the marginal contribution of each modality block, the following model configurations are evaluated:

Model	Blocks Included	Purpose
M1	A only	Financial baseline
M2	A + B (pooled)	Naive NLP augmentation
M2s	A + B (section-split)	Section-aware NLP
M3	A + B (section-split) + C	Full multimodal
M3e	A + scalar centrality + C	GNN vs. scalar ablation

Table 17: Ablation ladder configuration.

B.8 Research Output Artefact

All empirical results are operationalised through the **Deal Intelligence Terminal**, an interactive full-stack research artefact built in Streamlit/Python. The terminal renders the ablation results, hypothesis test statistics, SHAP feature attributions, and deal-level prediction scores dynamically, ensuring full reproducibility and providing a production-grade demonstration of the model’s practical utility.

B.9 Key Divergences from Initial Proposal

Initial Proposal	Final Implementation	Reason for Change
Integrate financial data + sentiment from news articles	Replaced news with pre-announcement 10-K filings	Filings are temporally precise, audited, and free from bias
Generic FinBERT sentiment	Section-split MDA vs. Risk Factor similarity	Discovered pooling cancels opposing signals (M2 reversal)
Crunchbase as data source	Replaced with Yahoo Finance + Bloomberg SPLC + SEC EDGAR	Crunchbase lacks coverage for pre-2015 institutional deals
ML framework for prediction accuracy	Extended to formal hypothesis testing + GNN architecture	Literature review revealed architectural limitations
Predict M&A success	Predict binary CAR direction	CAR is less behaviourally contaminated than accounting measures

Table 18: Summary of project trajectory and methodological divergences.

B.10 Ethical, Legal, and Social Considerations

All datasets used are publicly available or accessed via institutional licences (Bloomberg Terminal, LSEG Refinitiv). No human participants are involved. SEC filings are public regulatory documents. Supply-chain data is aggregated and does not include proprietary commercial relationships. Model outputs are not used for live trading or commercial decision-making within the scope of this study.

Submitted in partial fulfilment of CN6000, BSc Hons Data Science and Artificial Intelligence, University of East London.

Appendix C

Application for Approval of Research Activities

This appendix reproduces the ethical approval documentation submitted to and approved by the University of East London prior to the commencement of data collection. All research activities described in this dissertation were conducted under the terms of this approval.

The study involves no human participants, no personally identifiable data, and no experimental interventions. All data sources — Yahoo Finance equity returns, SEC EDGAR regulatory filings, and Bloomberg SPLC supply-chain records — are either publicly available or accessed through institutional licences held by the University. The ethical approval form confirms compliance with UEL's research governance framework and the British Psychological Society's Code of Ethics and Conduct.

The following pages contain the complete ethical approval application as submitted.

11/26/25, 11:55 AM

CN6000 Internal Ethical Approval Process 2025

The screenshot shows a web form titled "CN6000 Internal Ethical Approval Process 2025". At the top left is a red circular logo with the word "ETHICS" in blue. Below the logo is a warning: "Only complete this form if you have already received approval for the Proposal. If you complete this form before getting the final approval of your proposal, this form will be deleted, and you will be required to do it again!". The main heading is "Do you need internal approval?". Below this is a paragraph: "This section aims to check if you need internal approval in the first place. If you are not applying field research, then you may not need internal or external approval." There are two questions. Question 1: "1. Have you received the final Proposal Approval? That is the proposal has been reviewed by the 'second supervisor' and it has been approved? *". It has three radio button options: "Yes" (selected), "No", and "Maybe". Question 2: "2. My study will not involve collecting any data from humans. It is either just based on published datasets that I did not create or there are no datasets at all part of my study. *". It has two radio button options: "True" (selected) and "False". At the bottom, there is a URL: "https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRi0Z2BGW1qS2Ljw3259NmqdTAKKknKJUMU8yVThTWVPaV0ZQWV..." and "1/5".

ETHICS

CN6000 Internal Ethical Approval Process 2025

Only complete this form if you have already received approval for the Proposal. If you complete this form before getting the final approval of your proposal, this form will be deleted, and you will be required to do it again!

Do you need internal approval?

This section aims to check if you need internal approval in the first place. If you are not applying field research, then you may not need internal or external approval.

1. Have you received the final Proposal Approval? That is the proposal has been reviewed by the 'second supervisor' and it has been approved? *

☒ Yes

☐ No

☐ Maybe

2. My study will not involve collecting any data from humans. It is either just based on published datasets that I did not create or there are no datasets at all part of my study. *

☒ True

☐ False

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRi0Z2BGW1qS2Ljw3259NmqdTAKKknKJUMU8yVThTWVPaV0ZQWV...> 1/5

Figure 12: Ethical approval form — page 1 of 5.

1/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

3. Please state your topic title: * 

Illuminating the black box: an explainable AI framework for brain tumor segmentation using volumetric data interpretation in 3D CNNs.

4. Please state your supervisor's name (Spell it correctly). * 

Mimoona Maqsood SHARIF

5. If you answered **Yes** to question 1 and **True** to Question 2 - You are not required to proceed. Just click NEXT then Submit without the need to answer the remaining questions. Otherwise, please proceed to section 2. * 

☒ True

☐ False

This section aims to see if we can approve your proposal internally. 

Please read the questions carefully. Feel free to discuss these questions with your supervisor. If you find out that you have to apply for approval by the Research Ethics Committee, please discuss with your supervisor if there is a way to modify your dissertation in a way that does not take away from the value of your work.

6. I intend to plan and conduct my study as a pilot test of the research only. 

☐ Yes

☐ No

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRI0Z2BGW1qS2Ljw3259NmqdTakKknKJUMU8yVThTWVRaV0ZQWV...> 2/5

Figure 13: Ethical approval form — page 2 of 5.

1/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

☐ Maybe

7. I intend to conduct my field research as pilot test on UEL students or staff only. ☐

☐ Yes

☐ No

☐ Maybe

8. I will debrief participants before the study and mention that their contribution is entirely voluntary. ☐

☐ Yes

☐ No

☐ Maybe

9. I will not be collecting any personal data such as names, addresses, telephone numbers, emails...etc. ☐

☐ Yes

☐ No

☐ Maybe

10. I will not be asking any sensitive questions that can cause distress - such as mental health related or emotional questions. ☐

☐ Yes

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRI0Z2BGW1qS2Ljw3259NmqdTakKknKJUMU8yVThTWVPaV0ZQWV...> 3/5

Figure 14: Ethical approval form — page 3 of 5.

1/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

☐ No

☐ Maybe

11. I will be getting approval of the format and content of my questions from my supervisor before I proceed with my field research. ☐

☐ Yes

☐ No

☐ Maybe

12. I will have a 'consent' form for participants to be part of their study and be informed of their rights. ☐

☐ Yes

☐ No

☐ Maybe

13. I will store the data I collect on UEL OneDrive. No other copy of it will be stored elsewhere. ☐

☐ Yes

☐ No

☐ Maybe

14. I will not use the data for any other intended purpose and delete it at the end of the term. ☐

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRi0Z2BGW1qS2Ljw3259NmqdTakKknKJUMU8yVThTWVRaV0ZQWV...> 4/5

Figure 15: Ethical approval form — page 4 of 5.

1/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

☐ Yes

☐ No

☐ Maybe

15. If you have answered No or Maybe to any of the questions above, you cannot proceed with internal Ethical approval. I would advise you to review those particular questions with your supervisor before proceeding. If you both agree that you want to keep these answers, then please indicate below that you are happy to proceed. 

☐ All my answers are Yes. I have completed the internal approval and promise to abide by it.

☐ Not all my answers are yes. I have spoken to my supervisor and we both agreed that I should apply for form B Ethical Approval process.

16. I understand that violating any of the terms listed above could result in an automatic fail and possibly academic misconduct proceedings. 

☐ I confirm

 Microsoft 365

This content is created by the owner of the form. The data you submit will be sent to the form owner. Microsoft is not responsible for the privacy or security practices of its customers, including those of this form owner. Never give out your password.

Microsoft Forms | AI-Powered surveys, quizzes and polls [Create my own form](#)

[Privacy and cookies](#) | [Terms of use](#)

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRi0Z2BGW1qS2Ljw3259NmgdTakKknKJUMU8yVTtTWVRaV0ZQWV...> 5/5

Figure 16: Ethical approval form — page 5 of 5.

Appendix D

Deal Intelligence Terminal

The empirical results of this dissertation are operationalised through the **Deal Intelligence Terminal**, a full-stack interactive research artefact built in Streamlit and Python. The terminal is designed to bridge the gap between high-dimensional machine learning outputs and actionable financial decision-making. This appendix documents the terminal’s interface architecture, the diagnostic interaction flow, and the supplementary research modules used to defend the study’s methodological choices.

All screenshots were captured from a live instance of the terminal running against the production dataset described in Chapter 3.

D.1 The Model Evidence Wall

Before performing deal-specific diagnostics, the terminal provides a global **Ablation Wall** to validate the predictive superiority of the multimodal architecture. This interface allows researchers to verify that the inclusion of topological (GraphSAGE) and semantic (FinBERT) features yields a statistically significant lift in AUC-ROC and F1-score over the financial-only baseline, confirming the *Topological Alpha* and *Semantic Divergence* hypotheses prior to any deal-level inspection.




Figure 17: The Ablation Wall interface comparing the M1 (financial-only) baseline against the full M3 (financial + text + graph) architecture across AUC-ROC, F1, precision, and recall.

D.2 Deal-Level Diagnostics

The primary interface of the terminal is the **Diagnostic Report**, which provides a “Glass Box” view into a specific M&A transaction. The following screenshots document the diagnostic flow for the Omnicare → Theorem Clinical acquisition.

D.2.1 The Synergy Forecast

The top-level verdict translates raw model logits into a calibrated probability gauge. Alongside the synergy probability, the terminal displays two **Macro Metrics** for the deal:

- **Network Alpha:** The acquirer’s betweenness centrality rank within the global supply-chain graph, which informs the information transparency dampening effect (Hypothesis 3).
- **Semantic Match:** The cosine similarity between acquirer and target MD&A sections, indicating strategic alignment (Hypothesis 2).




Figure 18: The diagnostic verdict showing a 91% synergy probability calibrated via the M3 ensemble, alongside betweenness centrality rank (30th percentile) and MD&A semantic match.

D.2.2 Feature Interaction Zones

The middle section of the diagnostic report visualises the internal representation space of the model, separating the *where* (topology) from the *why* (semantics).




Figure 19: Zone 1 (Topological Embeddings) rendering the acquirer–target supply-chain sub-graph, and Zone 2 (Semantic Radar) plotting the pentagonal linguistic profile comparison.

- **Zone 1 — Topological Embeddings:** An interactive force-directed graph rendering the acquirer and target nodes within the global supply-chain ecosystem. This allows the user to visually inspect the degree of supplier overlap or customer risk that the GNN is encoding.
- **Zone 2 — Semantic Radar:** A pentagonal radar chart plotting the acquirer’s linguistic profile (MD&A, risk, strategic, sentiment, operational) against the target’s. Significant divergence in the Risk axis compared to the Strategic axis provides early warning signals of negative synergy.

D.2.3 The Glass Box (SHAP Attribution)

The final section of the diagnostic report addresses the “black box” criticism of GNNs and transformers by applying SHAP (SHapley Additive exPlanations) decomposition.




Figure 20: Zone 3: SHAP waterfall decomposition of the synergy probability, with the Algorithmic Translation summary below.

The waterfall plot decomposes the synergy probability into its constituent feature drivers. In this deal, the model identifies **Acquirer Operating Margin** and **Target Net Income** as the primary positive drivers.

The terminal also includes an **Algorithmic Translation** block — an NLP-driven summary that converts the mathematical SHAP variances into a natural language executive summary. This ensures that the *why* of the prediction is accessible to non-technical stakeholders without sacrificing the underlying econometric rigour.

D.3 Supplementary Research Modules

Beyond the primary diagnostic report, the terminal provides three specialised modules for methodological defence and hypothesis testing.

D.3.1 The Methodology Engine



Figure 21: The Methodology Engine interface, visualising the end-to-end M3 pipeline as an interactive Directed Acyclic Graph (DAG) with stage-specific output statistics.

The Methodology Engine serves as a live DAG of the study’s pipeline. It allows inspection of stage-specific output statistics (e.g., number of raw deals loaded versus text-coverage filtered deals) and provides direct links to the underlying Python implementation for each transformation step.

D.3.2 The Evaluation Lab



Figure 22: The Evaluation Lab, providing a structured defence of AUC-ROC as the primary performance metric with an interactive ROC curve visualisation.

The Evaluation Lab is designed to address the **Accuracy Paradox** inherent in imbalanced M&A datasets. It demonstrates that while a degenerate classifier could achieve 56% accuracy by predicting “Value Destruction” for all cases, the M3 architecture’s AUC of 0.5655 represents genuine predictive signal across threshold-independent decision boundaries.

D.3.3 The Hypothesis Lab




Figure 23: The Hypothesis Lab, showcasing the sector-wise AUC lift used to validate Hypothesis 1 (Topological Alpha) with $\pm 1\sigma$ error bars.

The Hypothesis Lab allows for real-time slicing of model performance across SIC-coded industries. This interface was used to generate the evidence for Hypothesis 1, demonstrating that the GraphSAGE modality provides a disproportionately higher lift in capital-heavy supply-chain sectors ($\Delta = +0.059$, $p = 0.0047$) compared to asset-light service industries ($\Delta = +0.041$, $p = 0.0329$).

D.3.4 Pipeline Architecture




Figure 24: The Pipeline Architecture interface, documenting the internal tensor flow and the 249-dimensional multimodal feature concatenation logic.

The final module of the terminal provides a high-level overview of the **Multimodal Fusion** head. It visualises the transformation of 67 financial features, 32 textual PCA dimensions, and 150 GraphSAGE embeddings into the final 249-dimensional feature vector consumed by the XGBoost inference engine.

Appendix E

Selected Implementation Code Snippets

This appendix presents selected implementation excerpts from the final system. The objective is to document the core engineering logic behind the dissertation’s empirical claims in a form that is readable, auditable, and aligned with the methodology described in Chapter 3.

E.1 Selection Criteria

Not every line of source code is reproduced here. The full project contains data ingestion scripts, model training utilities, evaluation pipelines, dashboard code, configuration files, and visualisation logic; reproducing the entire repository would obscure the methodological contribution rather than clarify it.

The selection criteria for this appendix are:

- Code that directly implements a methodological claim made in the dissertation;
- Code that protects against leakage or invalid evaluation;
- Code that defines a core modelling component;
- Code that supports interpretability or reproducibility.

Accordingly, the appendix focuses on six components: (1) event-study CAR generation, (2) temporal splitting with embargo control, (3) section-aware textual feature extraction, (4) heterogeneous graph embedding construction, (5) multimodal late-fusion prediction, and (6) SHAP-based interpretability.

E.2 Two-Stage CAR Pipeline

The first critical implementation block computes Cumulative Abnormal Return (CAR) from market data before any prediction model is trained. This separation matters because the dissertation explicitly decouples label construction from the downstream learning pipeline.

Purpose:

- Stage 1: estimate realised CAR using an OLS market model over the estimation window;
- Stage 2: train the prediction model using only pre-announcement features.

```
from scipy.stats import linregress
import numpy as np
import pandas as pd

ESTIMATION_START = -200
ESTIMATION_END = -20
EVENT_START = -5
EVENT_END = 5
MIN_OBS = 120

def compute_car_for_deal(df):
    est = df[(df["relday"] >= ESTIMATION_START) & (df["relday"] <=
ESTIMATION_END)]
    evt = df[(df["relday"] >= EVENT_START) & (df["relday"] <= EVENT_END)]

    est = est.dropna(subset=["acq_ret", "mkt_ret"])
    evt = evt.dropna(subset=["acq_ret", "mkt_ret"])

    if len(est) < MIN_OBS or evt.empty:
        return np.nan

    alpha, beta, _, _, _ = linregress(est["mkt_ret"], est["acq_ret"])
    evt = evt.copy()
    evt["expected_ret"] = alpha + beta * evt["mkt_ret"]
    evt["abnormal_ret"] = evt["acq_ret"] - evt["expected_ret"]

    return evt["abnormal_ret"].sum()

final_df["carm5p5"] = final_df.groupby("deal_id").apply(compute_car_for_deal)
final_df["car_binary"] = (final_df["carm5p5"] > 0).astype(int)
```

This excerpt shows that the target variable is generated from a conventional event-study process before the classifier or regressor is trained. It also makes clear that the binary target used in the classification pipeline is derived directly from the sign of CAR, not from post-hoc labelling.

E.3 Temporal Splitting and Embargo Control

A core validity claim in the dissertation is that the pipeline prevents forward-looking leakage through strict temporal splitting and an 11-trading-day embargo between partitions.

Purpose:

- Preserve chronological realism;
- Prevent overlapping CAR event windows across train, validation, and test sets;
- Ensure preprocessing is fitted only on training data.

```
import pandas as pd

EMBARGO_DAYS = 11

def apply_temporal_split(df):
    train = df[df["announce_year"].between(2000, 2016)].copy()
    valid = df[df["announce_year"].between(2017, 2019)].copy()
    test = df[df["announce_year"].between(2020, 2023)].copy()
    return train, valid, test

def embargo_boundary(left_df, right_df, left_max_date, right_min_date):
    left_keep = left_df[left_df["announce_date"] < left_max_date -
pd.Timedelta(days=EMBARGO_DAYS)]
    right_keep = right_df[right_df["announce_date"] > right_min_date +
pd.Timedelta(days=EMBARGO_DAYS)]
    return left_keep, right_keep

train, valid, test = apply_temporal_split(df)
train, valid = embargo_boundary(train, valid, train["announce_date"].max(),
valid["announce_date"].min())
valid, test = embargo_boundary(valid, test, valid["announce_date"].max(),
test["announce_date"].min())
```

This code operationalises one of the dissertation's strongest methodological safeguards. Without this step, neighbouring deals can share overlapping return windows and silently contaminate evaluation metrics.

E.4 Section-Aware FinBERT Extraction

The text pipeline deliberately separates Management Discussion and Analysis (MDA) from Risk Factors rather than treating the filing as a single undifferentiated document.

Purpose:

- Preserve section-specific semantics;
- Avoid signal cancellation between strategic-fit language and liability language;
- Compress document embeddings into a stable downstream representation.

```
from transformers import AutoTokenizer, AutoModel
import torch
import numpy as np

MODEL_NAME = "ProsusAI/finbert"
tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)
model = AutoModel.from_pretrained(MODEL_NAME)
model.eval()

def chunk_text(text, max_len=512, stride=256):
    tokens = tokenizer.encode(text, add_special_tokens=False)
    windows = []
    for i in range(0, max(1, len(tokens)), stride):
        chunk = tokens[i:i + max_len - 2]
        if not chunk:
            break
        windows.append([tokenizer.cls_token_id] + chunk +
[tokenizer.sep_token_id])
        if i + max_len - 2 >= len(tokens):
            break
    return windows

def embed_section(text):
    windows = chunk_text(text)
    vecs = []
    with torch.no_grad():
        for win in windows:
            x = torch.tensor([win])
            out = model(x, output_hidden_states=True)
            cls_vec = out.hidden_states[-2][0, 0, :].cpu().numpy()
            vecs.append(cls_vec)
    return np.mean(vecs, axis=0)

acq_mda_vec = embed_section(acq_item7_text)
acq_rf_vec = embed_section(acq_item1a_text)
```

This excerpt captures the exact modelling principle behind the Semantic Divergence Hypothesis: the two filing sections are processed independently and only merged later, rather than collapsed at source.

E.5 Heterogeneous Graph Construction and Embedding

The graph stream is the architectural feature that differentiates the final system from standard tabular M&A prediction pipelines.

Purpose:

- Encode firm position inside a supply-chain network;
- Distinguish upstream and downstream relationships through typed edges;
- Extract firm-level embeddings for downstream fusion.

```
from torch_geometric.data import HeteroData
from torch_geometric.nn import HeteroConv, SAGEConv
import torch.nn.functional as F
import torch

def build_hetero_graph(node_x, supplies_edges, buysfrom_edges):
    data = HeteroData()
    data["company"].x = node_x
    data[("company", "supplies", "company")].edge_index = supplies_edges
    data[("company", "buysfrom", "company")].edge_index = buysfrom_edges
    return data

class HeteroGraphSAGE(torch.nn.Module):
    def __init__(self, in_dim, hidden_dim=128, out_dim=64):
        super().__init__()
        self.conv1 = HeteroConv({
            ("company", "supplies", "company"): SAGEConv(in_dim, hidden_dim),
            ("company", "buysfrom", "company"): SAGEConv(hidden_dim, out_dim),
        }, aggr="mean")
        self.conv2 = HeteroConv({
            ("company", "supplies", "company"): SAGEConv(hidden_dim, out_dim),
            ("company", "buysfrom", "company"): SAGEConv(hidden_dim, out_dim),
        }, aggr="mean")

    def forward(self, x_dict, edge_index_dict):
        x_dict = self.conv1(x_dict, edge_index_dict)
        x_dict = {k: F.relu(v) for k, v in x_dict.items()}
        x_dict = self.conv2(x_dict, edge_index_dict)
        return x_dict["company"]
```

This is the clearest code-level expression of the Topological Alpha argument. Separate message-passing weights are learned for different relationship types, which would be impossible in a flat financial table.

E.6 Late-Fusion Prediction Head

The final predictive system uses late fusion rather than end-to-end joint training of all modalities.

Purpose:

- Keep each modality independently encoded;
- Permit controlled ablations by enabling or disabling individual streams;
- Produce a single fused representation for classification or regression.

```
import torch
import torch.nn as nn

class ProjectionHead(nn.Module):
    def __init__(self, in_dim, out_dim):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(in_dim, out_dim),
            nn.ReLU(),
        )

    def forward(self, x):
        return self.net(x)

class FusionModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.fin_proj = ProjectionHead(56, 64)
        self.txt_proj = ProjectionHead(128, 64)
        self.gph_proj = ProjectionHead(64, 32)
        self.classifier = nn.Sequential(
            nn.Linear(160, 64),
            nn.ReLU(),
            nn.Linear(64, 1),
        )

    def forward(self, fin_x, txt_x, gph_x):
        z_fin = self.fin_proj(fin_x)
        z_txt = self.txt_proj(txt_x)
        z_gph = self.gph_proj(gph_x)
        z = torch.cat([z_fin, z_txt, z_gph], dim=1)
        return self.classifier(z)
```

This excerpt shows the central design decision of the dissertation: modality-specific extraction first, multimodal combination second. It is also the component that makes the ablation ladder methodologically clean — disabling any one `ProjectionHead` produces a valid sub-model for comparison.

E.7 SHAP-Based Interpretation

The dissertation does not treat interpretability as optional presentation. SHAP is used to test whether the multimodal lift is economically credible at the feature level.

Purpose:

- Quantify feature-level contribution to prediction output;
- Verify that graph and text signals are genuinely used by the classifier;
- Generate visual evidence for the SHAP appendix figures and dashboard.

```
import shap
import numpy as np
import pandas as pd
from xgboost import XGBClassifier

model = XGBClassifier(eval_metric="auc", random_state=42)
model.fit(X_train, y_train)

explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_train)
mean_abs = np.abs(shap_values).mean(axis=0)

importance = pd.DataFrame({
    "feature": X_train.columns,
    "mean_abs_shap": mean_abs,
}).sort_values("mean_abs_shap", ascending=False)

importance.head(15)
```

This is the implementation bridge between the predictive model and the interpretability claims made in Chapter 4. It demonstrates how the study moves from raw classifier output to ranked economic drivers, ensuring that the SHAP attributions presented in the terminal and in the main body are computationally reproducible.

E.8 Repository Availability

The code excerpts in this appendix are intentionally selective. Full implementation files, configuration artefacts, and dashboard code are maintained in the project repository and can be supplied separately if examiners require full computational auditability.